

Winning Space Race with Data Science

Aliaksandr Karpovich
12/10/2023



Outline

- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

Executive Summary

- Summary of methodologies
 - Collecting data via REST API and Web Scrapping
 - Exploratory Data Analysis through visualisation
 - EDA with SQL
 - Interact with Map, Markers by using Folium
 - Interactive controls with Plotly Dash
 - Predictive Analysis
- Summary of all results
 - EDA observations
 - Geo-data observations
 - Graphs with interactive controls
 - Predicted results

Introduction

- Project background and context
 - The idea of the project is to predict outcome of Falcon 9 landing. Since launch of the Falcon 9 rocket costs 62 million dollars it's reasonable to use different attributes and params to predict its landing and by doing this determine the cost of a launch.
- Problems you want to find answers
 - What attributes make landing successful?
 - How each of the attributes affect on landing result?
 - What launch sites allow SpaceX to achieve the best success rate?

Section 1

Methodology

Methodology

Executive Summary

- Data collection methodology:
 - SpaceX REST API
 - Web Scrapping from Wikipedia

```
static_json_url='https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork'
```

We should see that the request was successfull with the 200 status response code

```
response.status_code
```

```
200
```

Now we decode the response content as a Json using `.json()` and turn it into a Pandas dataframe using `.json_normalize()`

```
# Use json_normalize meethod to convert the json result into a dataframe  
data = pd.json_normalize(response.json())
```

```
# use requests.get() method with the provided static_url  
# assign the response to a object  
response = requests.get(static_url)
```

Create a `BeautifulSoup` object from the HTML `response`

```
# Use BeautifulSoup() to create a BeautifulSoup object from a response text content  
soup = BeautifulSoup(response.text)
```

Print the page title to verify if the `BeautifulSoup` object was created properly

```
# Use soup.title attribute  
soup.title
```

Methodology

Executive Summary

- Perform data wrangling

```
# HINT: Use get_dummies() function on the categorical columns
features_one_hot = pd.get_dummies(df)
print(features_one_hot)
```

	FlightNumber	PayloadMass	Flights	GridFins	Reused	Legs	Block	\
0	1	6104.959412	1	False	False	False	1.0	
1	2	525.000000	1	False	False	False	1.0	
2	3	677.000000	1	False	False	False	1.0	
3	4	500.000000	1	False	False	False	1.0	
4	5	3170.000000	1	False	False	False	1.0	
..
85	86	15400.000000	2	True	True	True	5.0	
86	87	15400.000000	3	True	True	True	5.0	
87	88	15400.000000	6	True	True	True	5.0	
88	89	15400.000000	3	True	True	True	5.0	
89	90	3681.000000	1	True	False	True	5.0	

Methodology

Executive Summary

- Perform exploratory data analysis (EDA) using visualization and SQL
- Perform interactive visual analytics using Folium and Plotly Dash
- Perform predictive analysis using classification models
 - How to build, tune, evaluate classification models

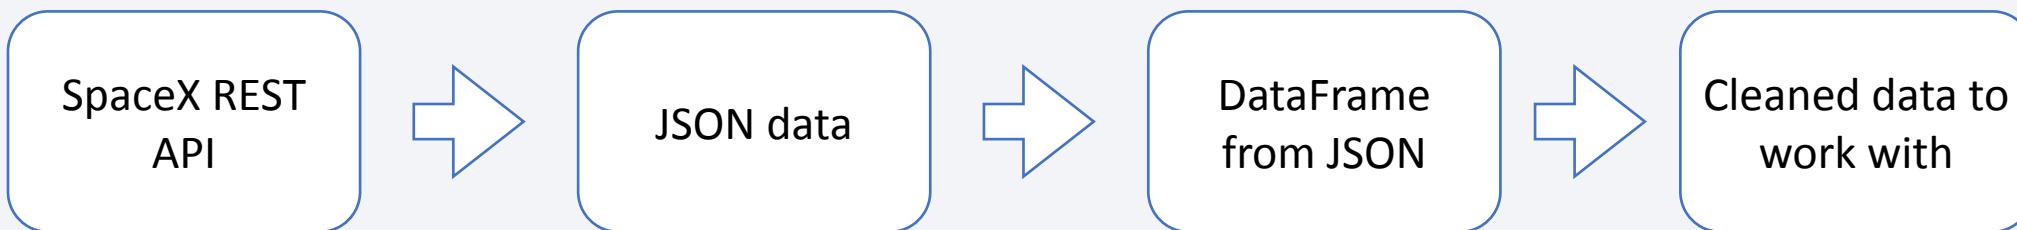
```
parameters ={'C':[0.01,0.1,1],  
            'penalty':['l2'],  
            'solver':['lbfgs']}
```

```
parameters ={"C": [0.01, 0.1, 1], 'penalty': ['l2'], 'solver': ['lbfgs']}# l1 lasso l2 ridge  
lr=LogisticRegression()  
logreg_cv = GridSearchCV(lr, parameters, cv=10)  
logreg_cv.fit(X_train, Y_train)
```

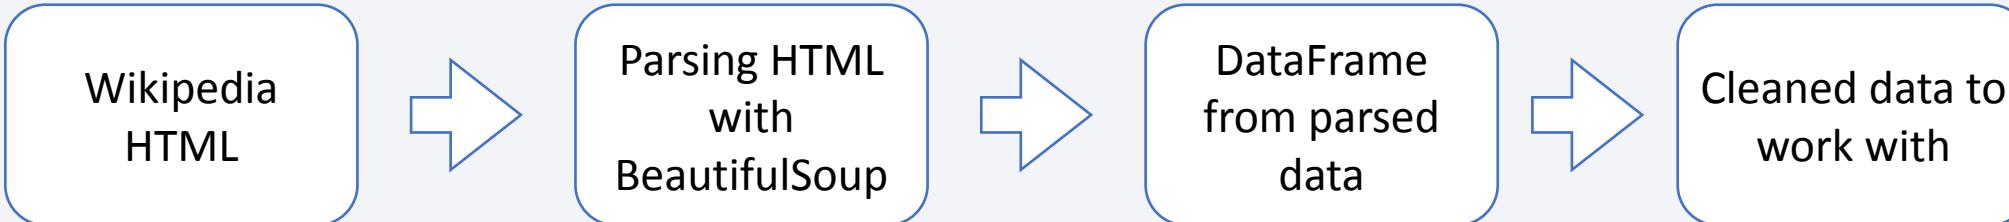
```
GridSearchCV(cv=10, estimator=LogisticRegression(),  
            param_grid={'C': [0.01, 0.1, 1], 'penalty': ['l2'],  
                        'solver': ['lbfgs']})
```

Data Collection

- The info obtained by the SpaceX REST API using the endpoint `api.spacexdata.com/v4/` is rocket, launches, payload mass.



- The info obtained from Wikipedia https://en.wikipedia.org/w/index.php?title=List_of_Falcon_9_and_Falcon_Heavy_launches&oldid=1027686922 by using web scrapping.



Data Collection – SpaceX API

Collecting data from SpaceX API:



[Link to Github with the notebook](#)

Data Collection - Scraping

Collecting Data by using web scrapping

```
static_url = "https://en.wikipedia.org/w/index.php?title=List_of_Falcon_9_and_Falcon_Heavy_launches&oldid=1027686922"
```

Download html content of the page
response = requests.get(static_url)

Parse the html using BeautifulSoup library
soup = BeautifulSoup(response.text)

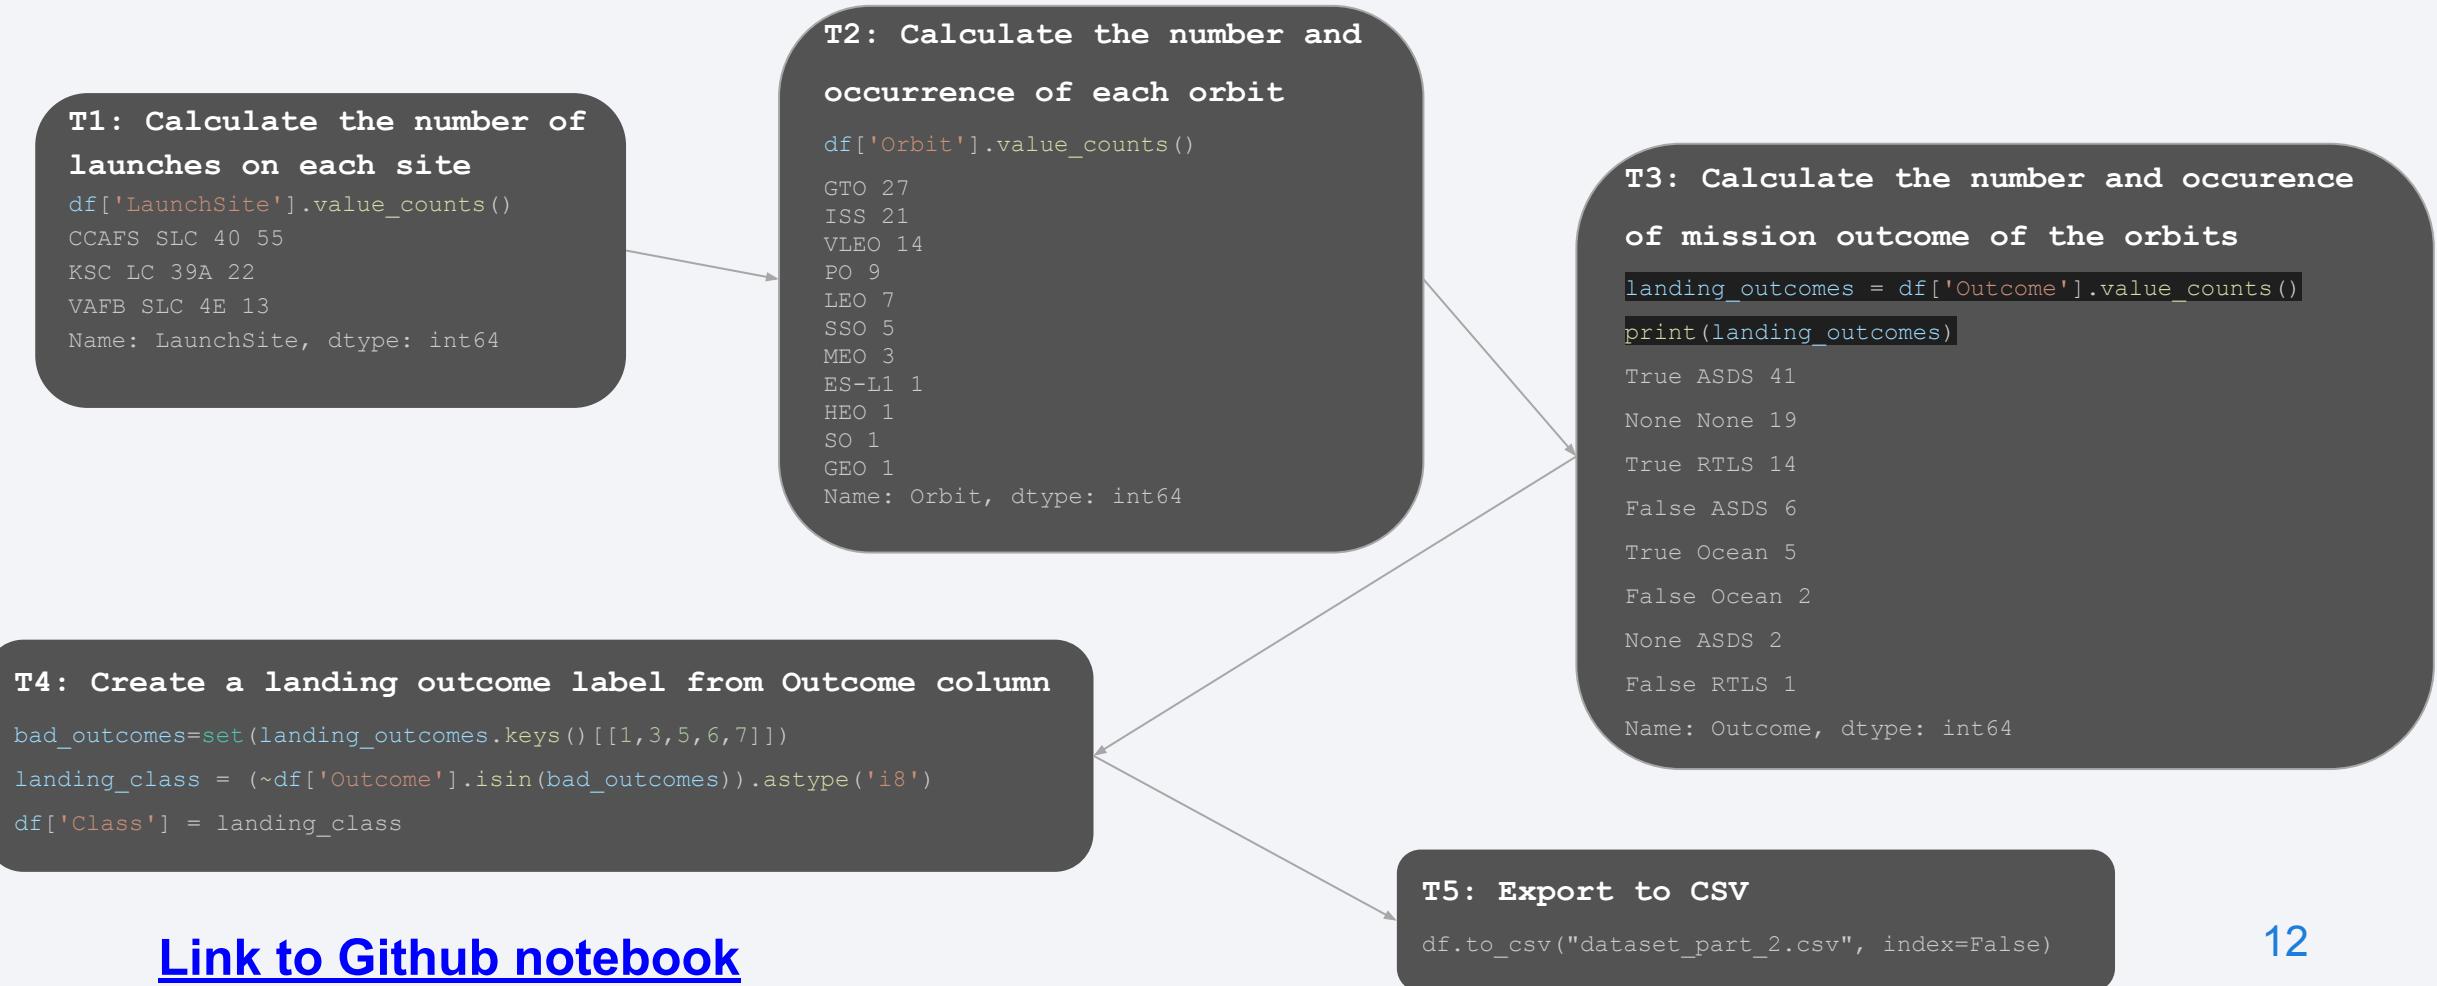
Create DataFrame from dict with parsed data
df= pd.DataFrame({ key:pd.Series(value) for key, value in launch_dict.items()})

Export data to csv
df.to_csv('spacex_web_scraped.csv', index=False)

[Link to Github with the notebook](#)

Data Wrangling

- The Outcome string variable has to be transformed into categorical variable which equals to 1 if land is successful and 0 otherwise.



[Link to Github notebook](#)

EDA with Data Visualization

Scatter Graphs

- Flight Number vs. Payload Mass
- Flight Number vs. Launch Site
- Payload vs. Launch Site
- Orbit vs. Flight Number
- Payload vs. Orbit Type
- Orbit vs. Payload Mass



shows the relationship between two quantitative variables measured for the same individuals

Bar Graph

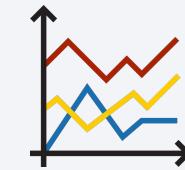
- Success rate vs. Orbit



shows the relationship between a numeric and a categorical variable

Line Graph

- Success rate vs. Year

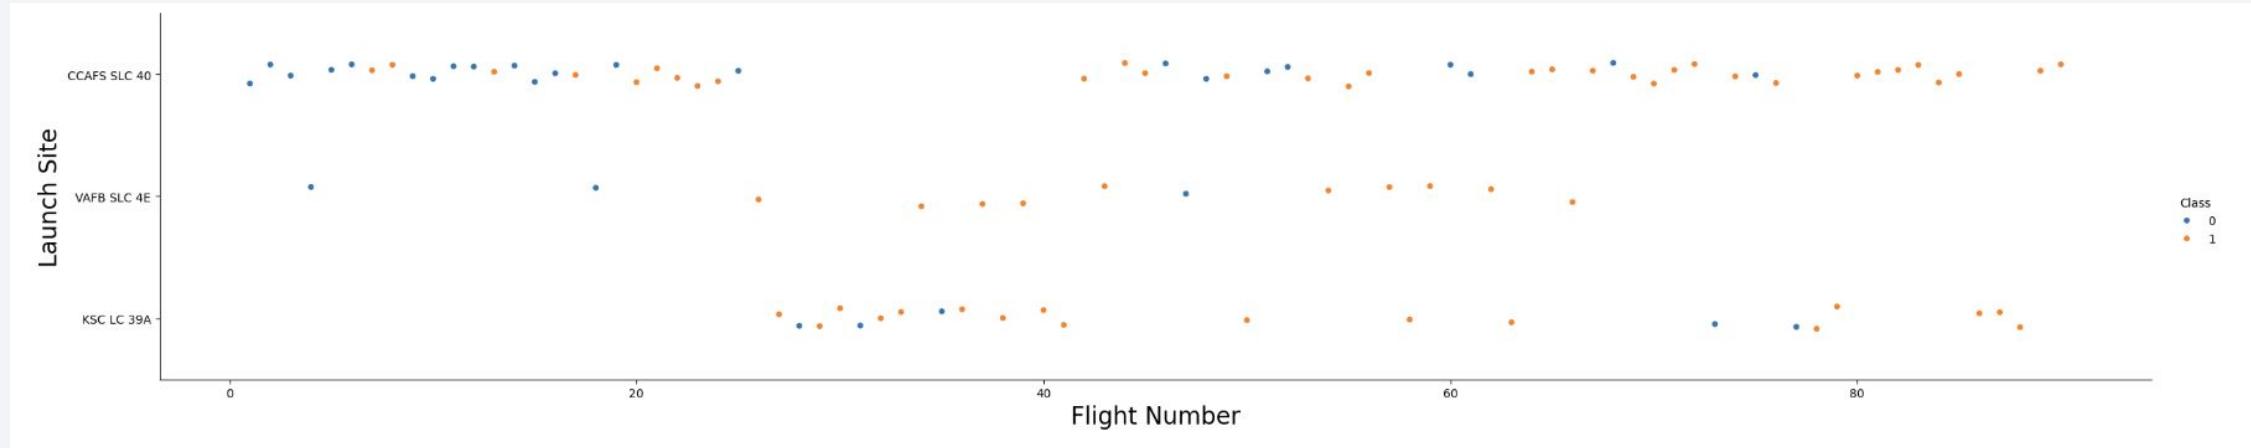


shows data along a number line with Xs or dots recorded above the responses to indicate the number of occurrences a response appears in the data set

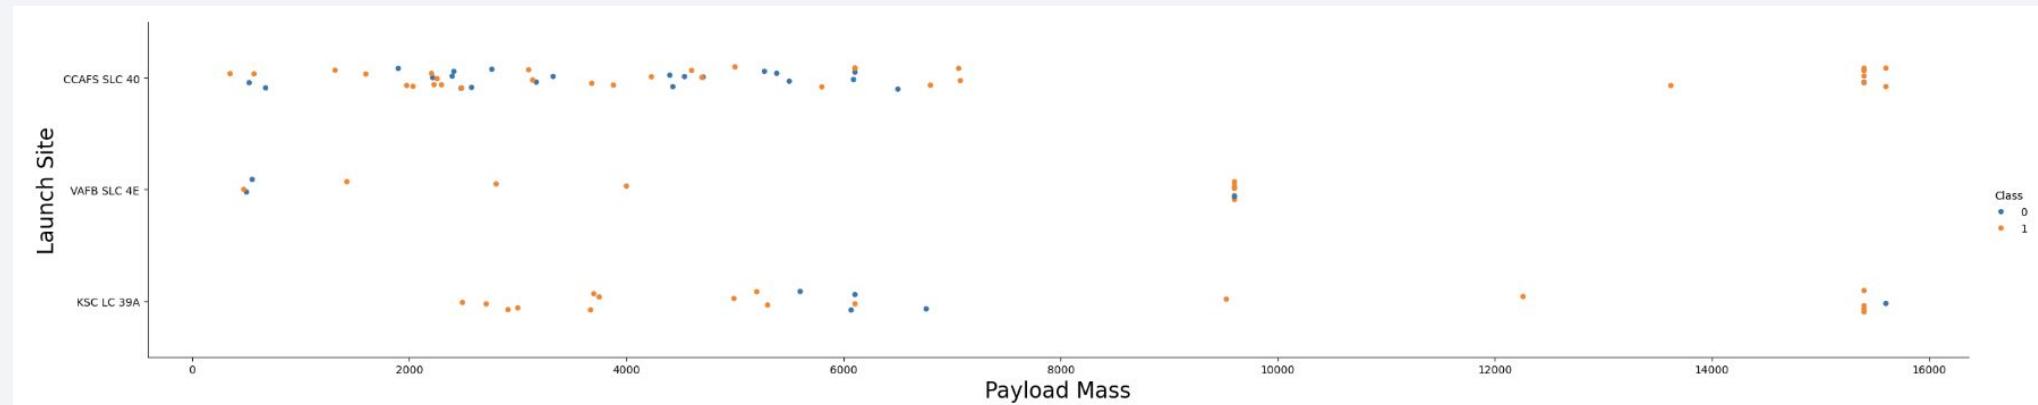
[Link to Github notebook](#)

EDA with Data Visualization

Flight Number vs Launch Site

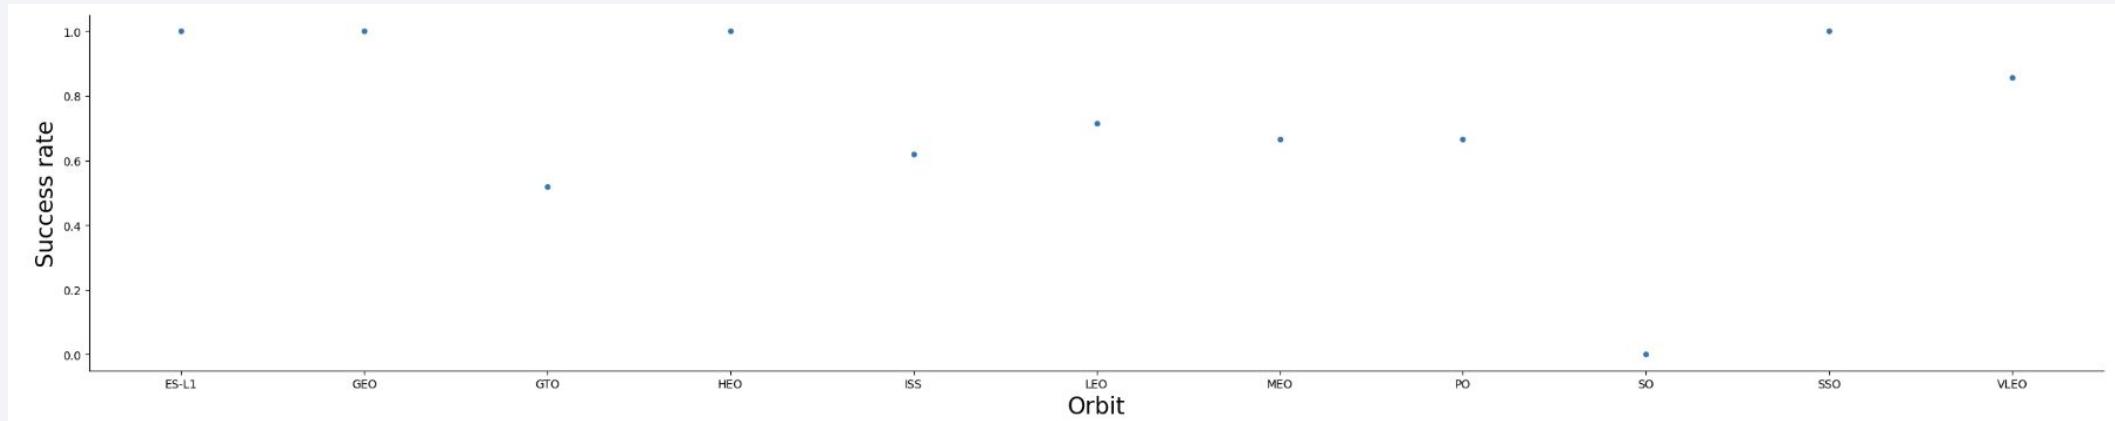


Payload vs Launch Site

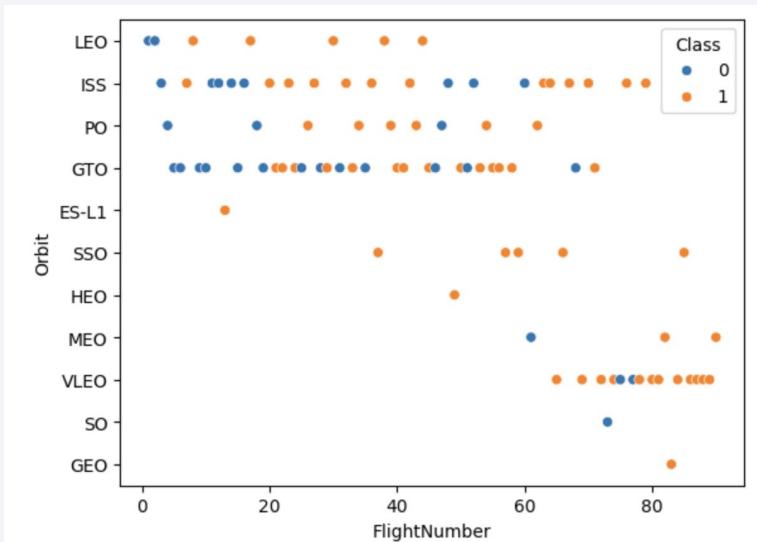


EDA with Data Visualization

Success rate vs Orbit

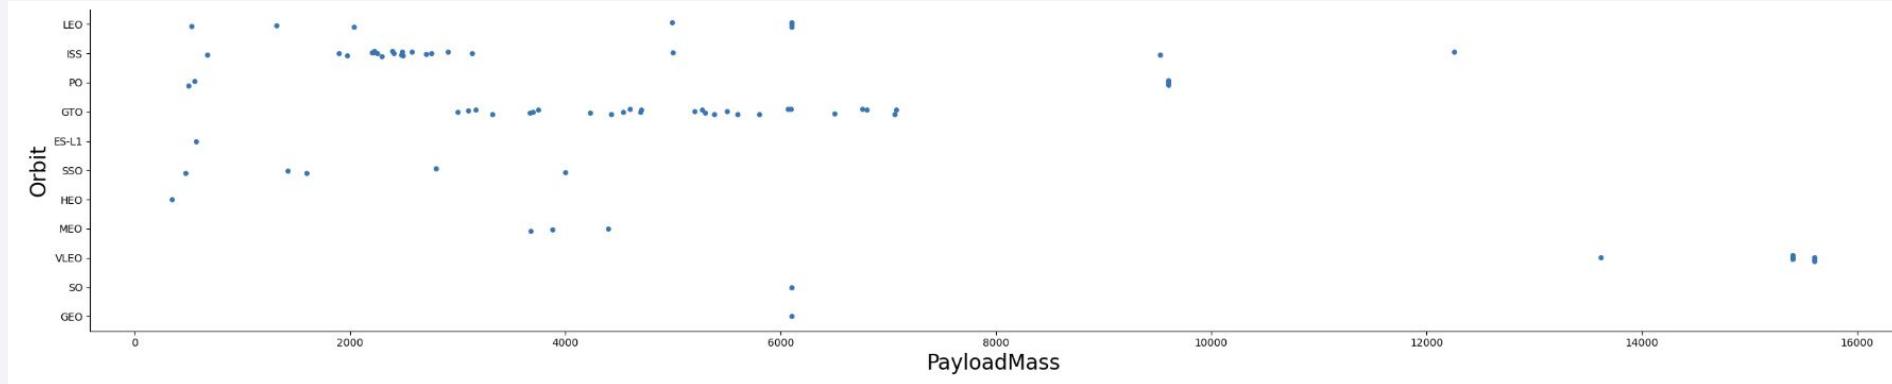


Flight Number vs Orbit

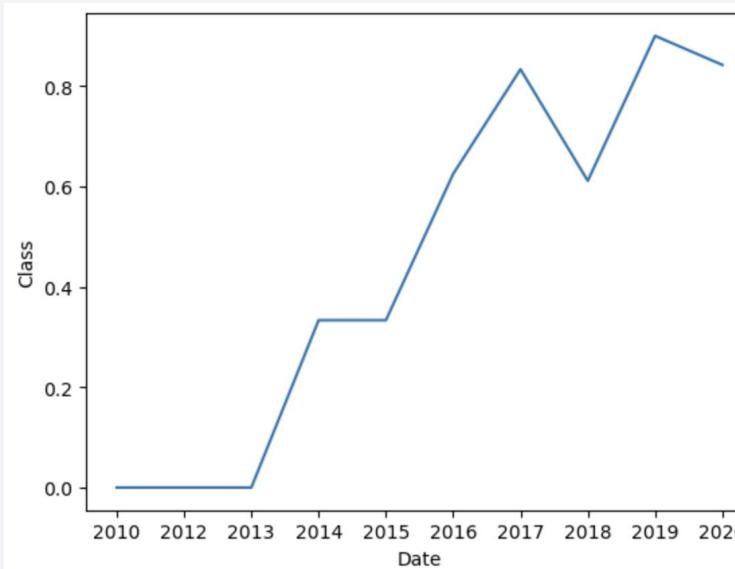


EDA with Data Visualization

Payload vs Orbit



Launch success yearly trend line



EDA with SQL

- We performed SQL queries to gather and understand data from dataset:
 - Displaying the names of the unique launch sites in the space mission.
 - Display 5 records where launch sites begin with the string 'CCA'
 - Display the total payload mass carried by boosters launched by NASA (CRS).
 - Display average payload mass carried by booster version F9 v1.1.
 - List the date when the first successful landing outcome in ground pad was achieved.
 - List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000.
 - List the total number of successful and failure mission outcomes.
 - List the names of the booster_versions which have carried the maximum payload mass.
 - List the records which will display the month names, failure landing_outcomes in drone ship ,booster versions, launch_site for the months in year 2015.
 - Rank the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order.

[Link to Github notebook](#)

All Launch Site Names

SQL Query:

```
select distinct launch_site from SPACEXTABLE
```

Launch_Site

CCAFS LC-40

VAFB SLC-4E

KSC LC-39A

CCAFS SLC-40

Explanation

It runs through all launch_site records and returns unique with DISTINCT operator.

Launch Site Names Begin with 'CCA'

```
%sql select * from SPACEXTABLE where launch_site like 'CCA%' limit 5
```

```
* sqlite:///my_data1.db
```

```
Done.
```

Date	Time (UTC)	Booster_Version	Launch_Site	Payload	PAYLOAD_MASS_KG_	Orbit	Customer	Mission_Outcome	Landing.
2010-04-06	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit	0	LEO	SpaceX	Success	Failure (1)
2010-08-12	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of Brouere cheese	0	LEO (ISS)	NASA (COTS) NRO	Success	Failure (1)
2012-05-22	07:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2	525	LEO (ISS)	NASA (COTS)	Success	N
2012-08-10	00:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1	500	LEO (ISS)	NASA (CRS)	Success	N
2013-01-03	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2	677	LEO (ISS)	NASA (CRS)	Success	N

Explanation

It filters records by launch_site to start with CCA and limits the output to 5 records

Total Payload Mass

```
%sql select SUM(PAYLOAD_MASS__KG_) from SPACEXTABLE where customer="NASA (CRS)"
```

```
* sqlite:///my_data1.db
```

```
Done.
```

SUM(PAYLOAD_MASS__KG_)
45596

Explanation:

It sums payload_mass of all NASA (CRS) customers.

Average Payload Mass by F9 v1.1

```
%sql select AVG(PAYLOAD_MASS_KG_) from SPACEXTABLE where booster_version='F9 v1.1'
```

```
* sqlite:///my_data1.db  
Done.
```

AVG(PAYLOAD_MASS_KG_)
2928.4

Explanation:

It filters records by booster_version to be 'F9 v1.1' and calculates average of payload_mass using built-in AVG function.

First Successful Ground Landing Date

```
%sql select min(date) from SPACEXTABLE where landing_outcome="Success (ground pad)"
```

```
* sqlite:///my_data1.db
```

```
Done.
```

min(date)
2015-12-22

Explanation:

It filters records by landing_outcome to be “Success (ground pad)” and finds minimum date using built-in MIN function.

Successful Drone Ship Landing with Payload between 4000 and 6000

```
%sql select Booster_Version from SPACEXTABLE where landing_outcome="Success (drone ship)" and PAYLOAD_MASS__KG_ between 4000 and 6000  
* sqlite:///my_data1.db  
  
Done.  
-----  
Booster_Version  
F9 FT B1022  
F9 FT B1026  
F9 FT B1021.2  
F9 FT B1031.2
```

Explanation:

It filters records by landing outcome to be “Success (drone ship)” and payload mass to be in range [4000, 6000] kg and returns only booster version.

Total Number of Successful and Failure Mission Outcomes

```
%sql select (select COUNT(*) from SPACEXTABLE where mission_outcome='Success') as success_count,  
(select COUNT(*) from SPACEXTABLE where mission_outcome!= 'Success') as failure_count
```

```
* sqlite:///my_data1.db
```

Done.

```
.....  
success_count  failure_count  
-----  
         98          3
```

Explanation:

It uses 2 subqueries, one of them counts all records with successful mission outcome and another with not successful.

Boosters Carried Maximum Payload

```
%sql select booster_version from SPACEXTABLE where PAYLOAD_MASS__KG_ = (select MAX(PAYLOAD_MASS__KG_) from SPACEXTABLE)
* sqlite:///my_data1.db
```

Done.

Booster_Version
F9 B5 B1048.4
F9 B5 B1049.4
F9 B5 B1051.3
F9 B5 B1056.4
F9 B5 B1048.5
F9 B5 B1051.4
F9 B5 B1049.5
F9 B5 B1060.2
F9 B5 B1058.3
F9 B5 B1051.6
F9 B5 B1060.3
F9 B5 B1049.7

Explanation:

At first it finds max payload mass using subquery and then returns records with such payload mass.

2015 Launch Records

```
%sql select substr(Date,6,2) as month, landing_outcome, booster_version, launch_site  
|from SPACEXTABLE where landing_outcome='Failure (drone ship)' and substr(Date,0,5) == '2015'  
* sqlite:///my_data1.db
```

Done.

.....

month	Landing_Outcome	Booster_Version	Launch_Site
10	Failure (drone ship)	F9 v1.1 B1012	CCAFS LC-40
04	Failure (drone ship)	F9 v1.1 B1015	CCAFS LC-40

Explanation:

It filters records by landing outcome to be Failure (drone ship) and by year to be 2015. To get year from a date it uses substr function.

Rank Landing Outcomes Between 2010-06-04 and 2017-03-20

```
%sql select landing_outcome, Count(landing_outcome) as cnt from SPACEXTABLE  
where landing_outcome in ("Failure (drone ship)", "Success (ground pad)") and  
date between '2010-06-04' and '2017-03-20'  
GROUP BY landing_outcome  
ORDER BY cnt desc
```

```
* sqlite:///my_data1.db  
Done.
```

Landing_Outcome	cnt
Success (ground pad)	5
Failure (drone ship)	5

Explanation:

It groups records filtered records by landing outcome and for each outcome counts its records

Build an Interactive Map with Folium

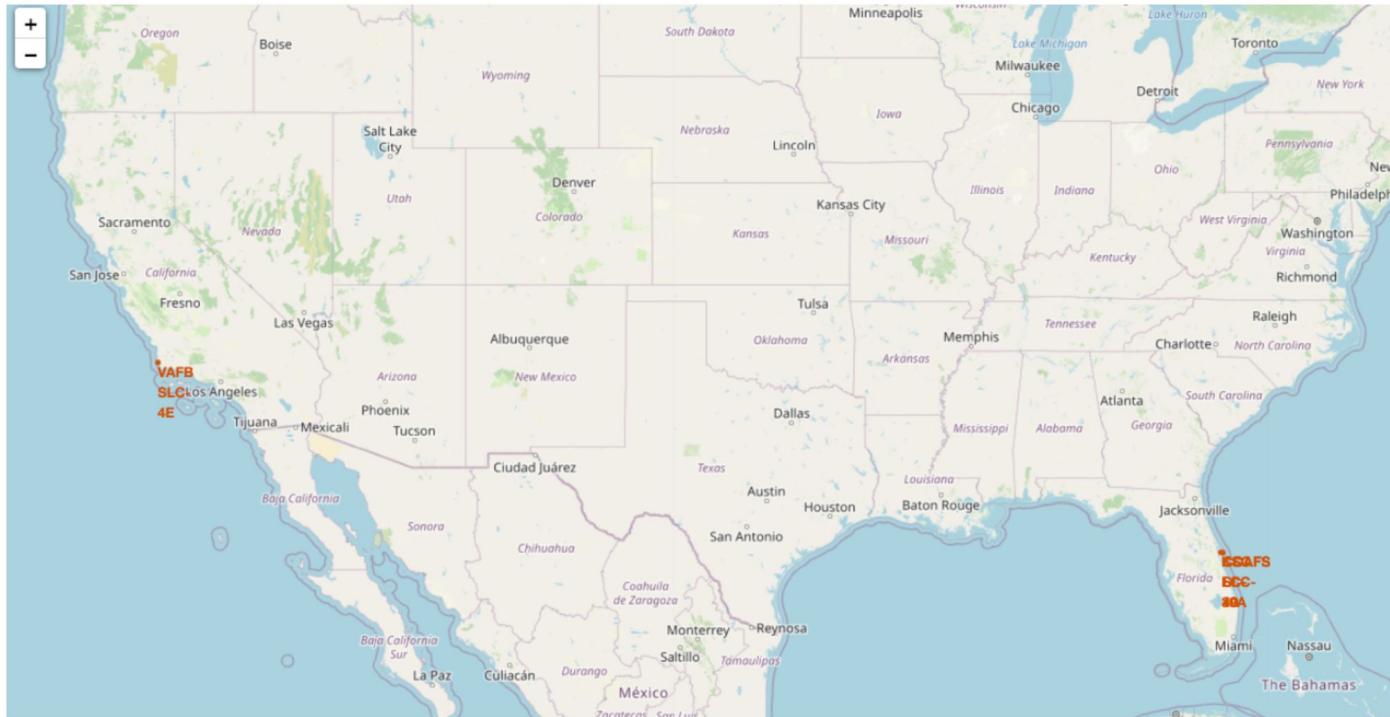
Completed tasks:

- Created a folium `Map` object, with an initial center location to be NASA Johnson Space Center at Houston, Texas.
`nasa_coordinate = [29.559684888503615, -95.0830971930759]`
`site_map = folium.Map(location=nasa_coordinate, zoom_start=10)`
- Created and added `folium.Circle` and `folium.Marker` for each launch site on the site map
- For each launch result in `spacex_df` data frame, added a `folium.Marker` to `marker_cluster`
- *Calculate the distances between a launch site to its proximities*
- *Added a MousePosition* on the map to get coordinate for a mouse over a point on the map
`formatter = "function(num) {return L.Util.formatNum(num, 5);}"`
`mouse_position = MousePosition(`
 `position='topright',`
 `separator=' Long: ',`
 `empty_string='NaN',`
 `long_first=False,`
 `num_digits=20,`
 `prefix='Lat: ',`
 `lat_formatter=formatter,`
 `long_formatter=formatter,`
`)`
`site_map.add_child(mouse_position)`
- *Mark down a point on the closest coastline using MousePosition and calculate the distance between the coastline point and the launch site*
- *Draw a PolyLine between a launch site to the selected coastline point*
- *Created a marker with distance to a closest railway and drew a line between the marker to the launch site*

It helps to answer such questions as:

- Are launch sites in close proximity to railways?
- Are launch sites in close proximity to highways?
- Are launch sites in close proximity to coastline?
- Do launch sites keep certain distance away from cities?

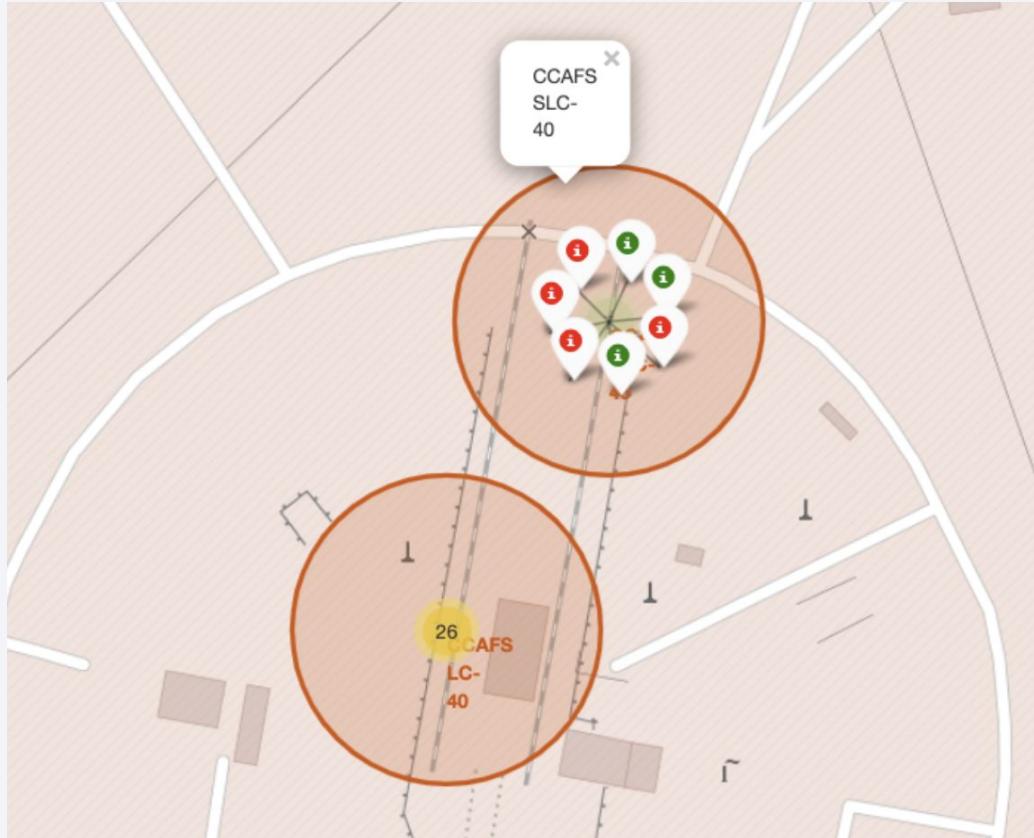
Folium Map - Ground stations



Explanation:

On that map we see that SpaceX launch sites located on the coast of US.

Folium Map - Marker clusters and labels

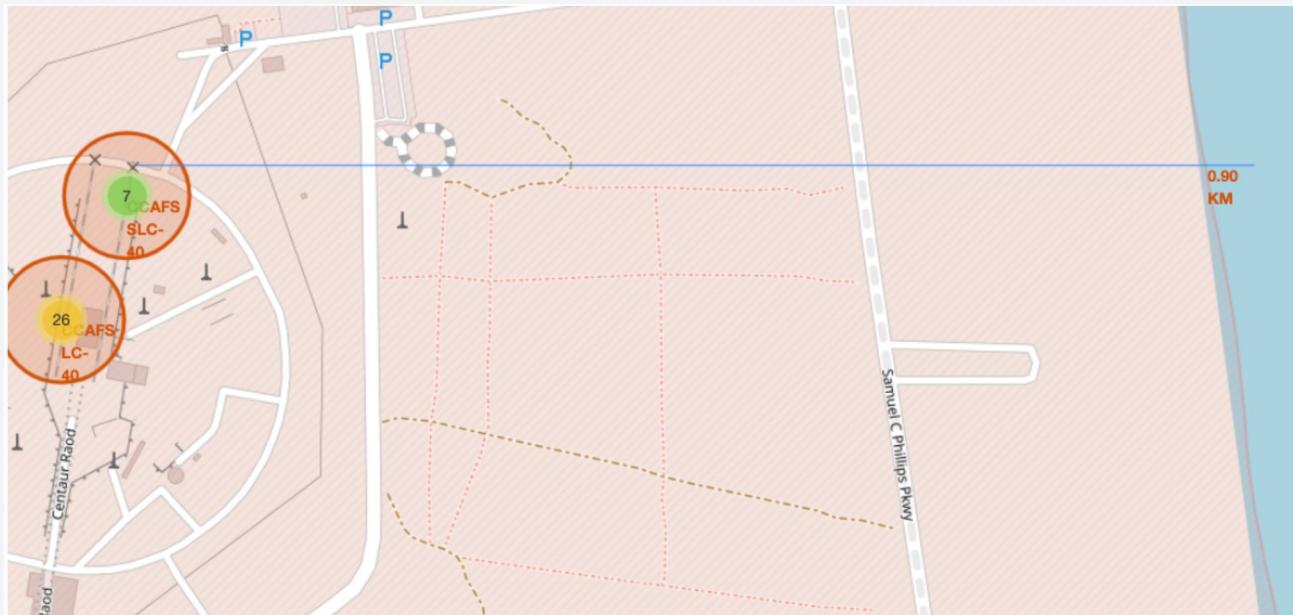


Explanation:

Red labels represent Failed launches and green labels represent Successful launches.

The most successful launch site is KSC LC-39A.

Folium Map - Distances between CCAFS SLC-40 and its proximities



Explanation:

CCAFS SLC-40 is pretty close to all of its proximities such as railways, highways, coastline but it doesn't keep distance away from cities.

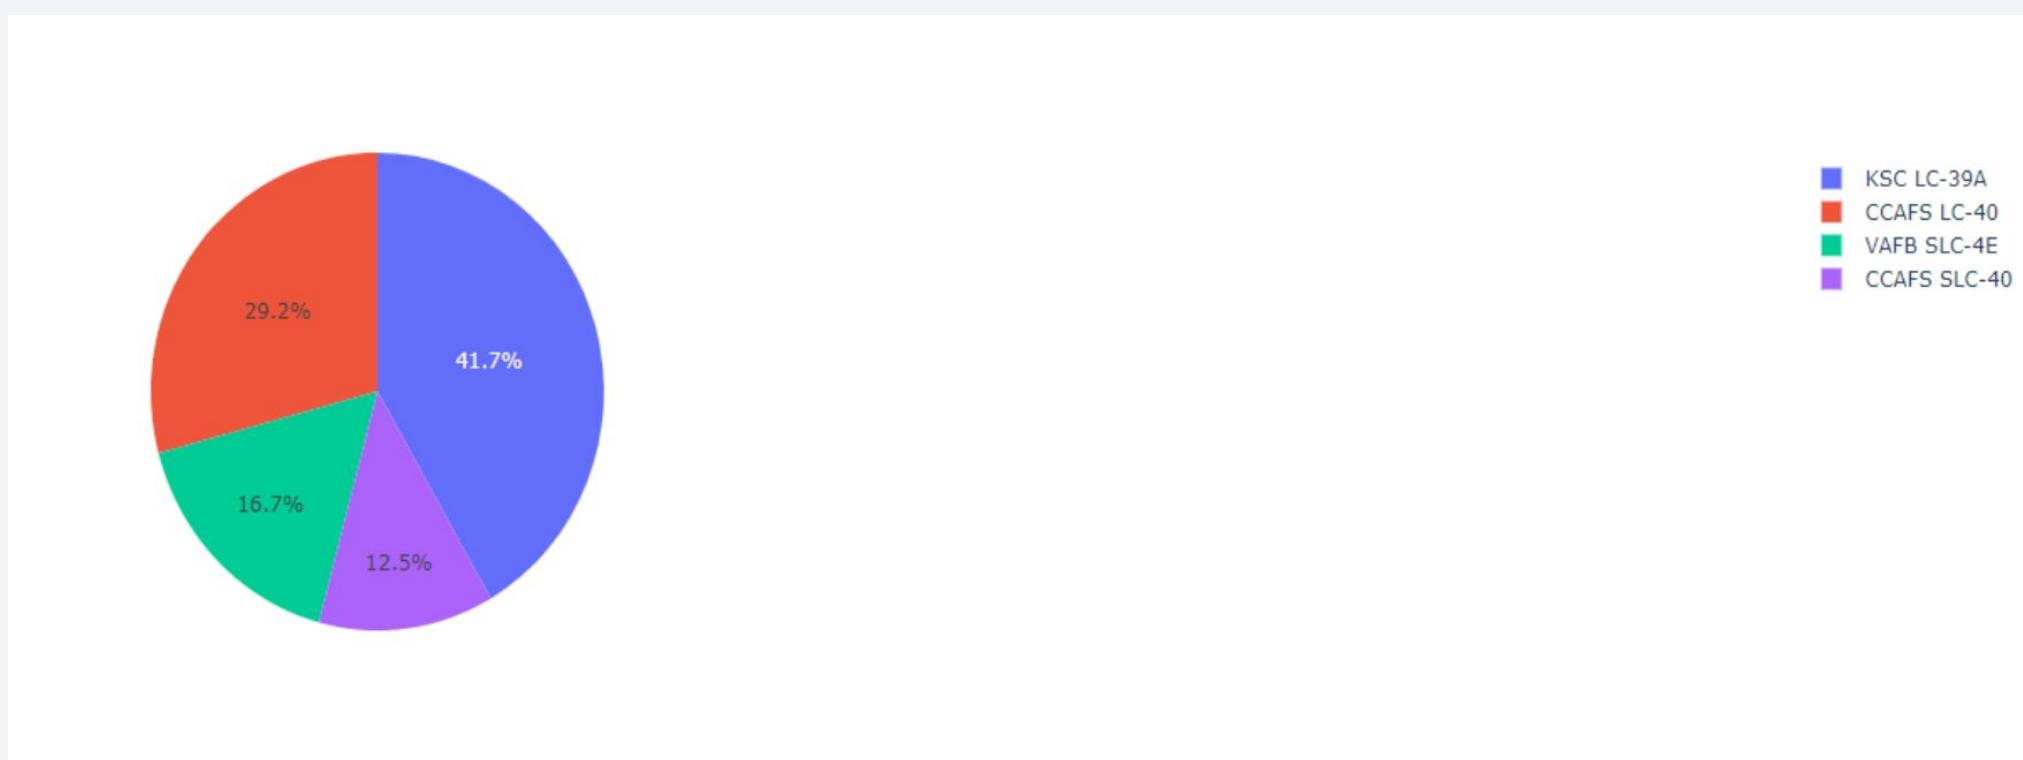
Build a Dashboard with Plotly Dash

Dashboard has dropdown, pie chart, rangeslider and scatter plot components:

- Dropdown allows a user to choose either specific launch site or ALL.
- Pie chart shows either success percentage among ALL launch sites or success/failure percentage for specific launch site.
- Rangeslider allows a user to control a payload mass range.
- Scatter chart shows the relationship between Success vs Payload Mass.

[Link to Github file](#)

Plotly Dash - Total success rate



Explanation:

KSL LC-39A has the highest success rate.

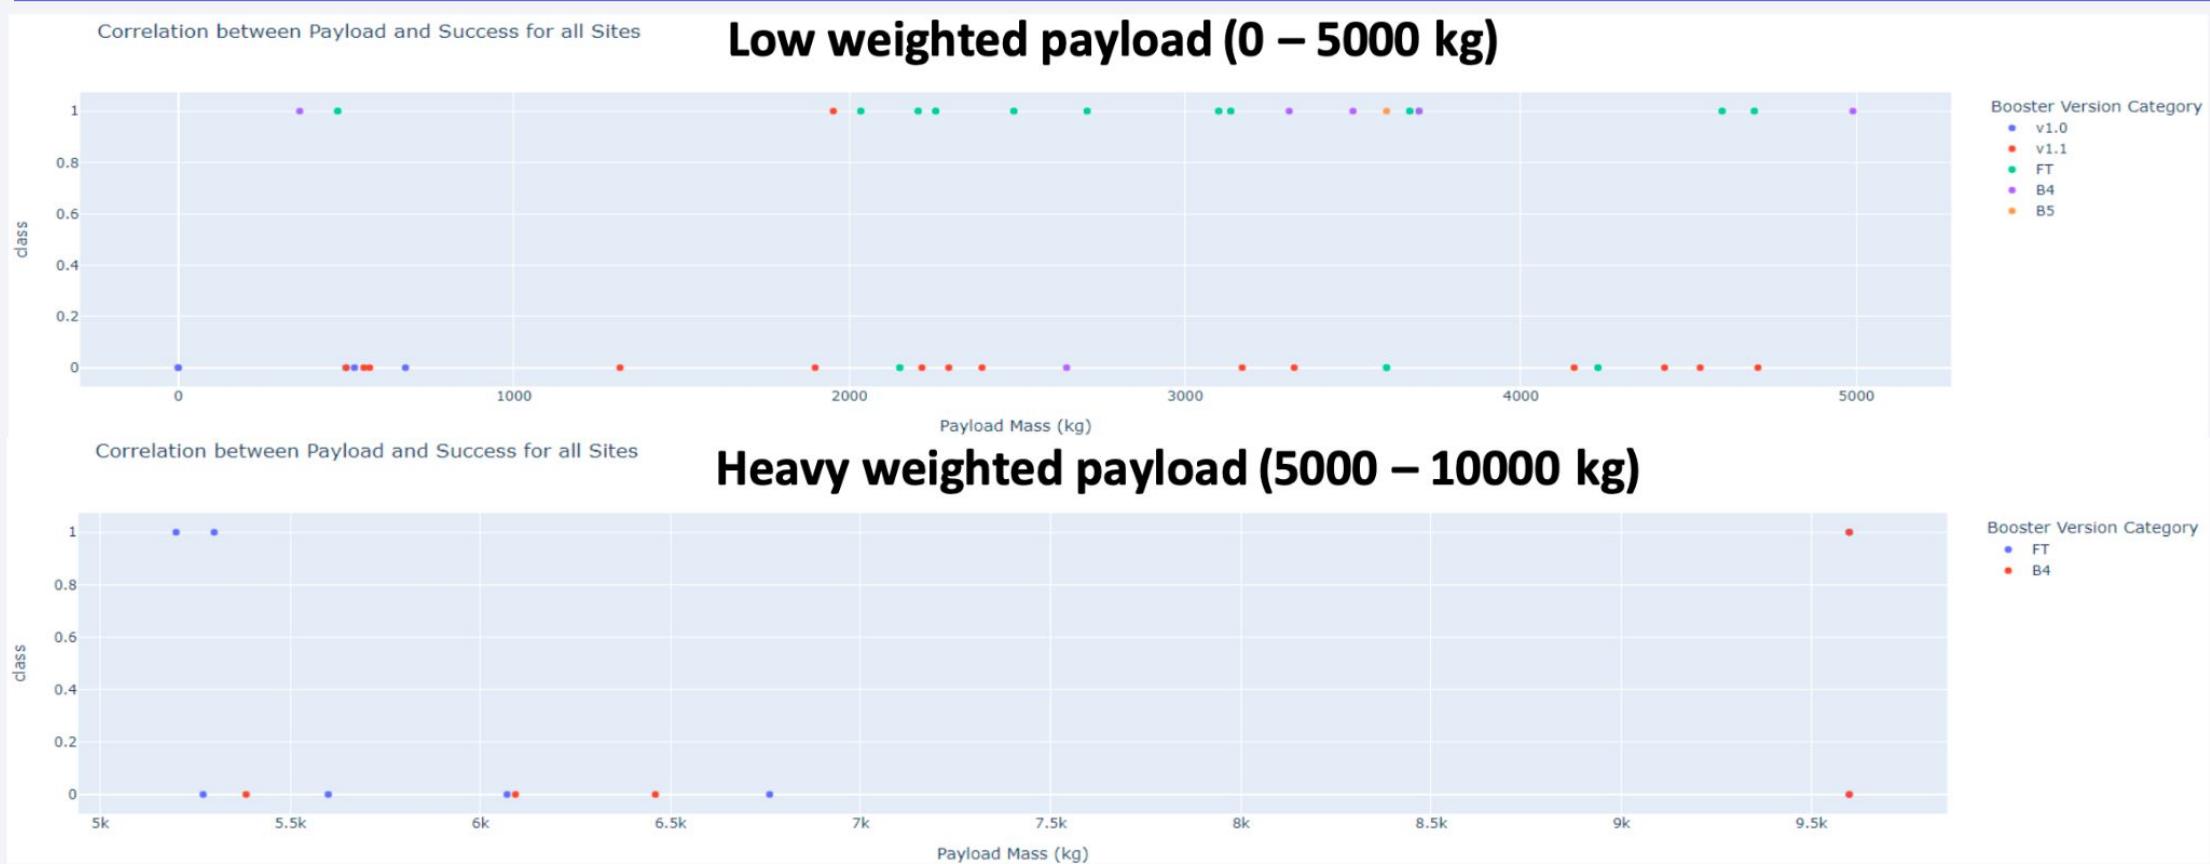
Plotly Dash - KSL LC-39A success rate



Explanation:

The launch site has 76.9% success launches and only 23.1% failed.

Plotly Dash - Correlation between Payload Mass vs Outcome for all launch sites



Explanation:

By comparing the plots we can see that success outcome for specific Booster versions is highest payload in [2000, 5000] kg.

Predictive Analysis (Classification)

Performed steps to find the best model and its params:

- Load data
- Pick X features and Y labels variables
- Standardize X variables
- Split data into train and test subsets
- Select the best hyperparameters to describe training data using GridSearchCV and the following models:
 - Logistic regression
 - Support vector machine
 - Decision tree classifier
 - K nearest neighbors
- Compare models on test data

[Link to Github notebook](#)

Predictive Analysis (Classification)

- Logistic Regression

```
parameters ={'C':[0.01,0.1,1],  
            'penalty':['l2'],  
            'solver':['lbfgs']}
```

```
parameters ={"C":[0.01,0.1,1], 'penalty':['l2'], 'solver':['lbfgs']}# l1 lasso l2 ridge  
lr=LogisticRegression()  
logreg_cv = GridSearchCV(lr, parameters, cv=10)  
logreg_cv.fit(X_train, Y_train)
```

```
GridSearchCV(cv=10, estimator=LogisticRegression(),  
            param_grid={'C': [0.01, 0.1, 1], 'penalty': ['l2'],  
                        'solver': ['lbfgs']})
```

```
print("tuned hpyerparameters :(best parameters) ",logreg_cv.best_params_)  
print("accuracy :",logreg_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters)  {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}  
accuracy : 0.8464285714285713
```

Predictive Analysis (Classification)

- Support Vector Machine

```
parameters = {'kernel':('linear', 'rbf','poly','rbf', 'sigmoid'),
              'C': np.logspace(-3, 3, 5),
              'gamma':np.logspace(-3, 3, 5)}
svm = SVC()

svm_cv = GridSearchCV(svm, parameters, cv=10)
svm_cv.fit(X_train, Y_train)

GridSearchCV(cv=10, estimator=SVC(),
             param_grid={'C': array([1.0000000e-03, 3.16227766e-02, 1.0000000e+00, 3.16227766e+01,
1.0000000e+03]),
              'gamma': array([1.0000000e-03, 3.16227766e-02, 1.0000000e+00, 3.16227766e+
01,
1.0000000e+03]),
              'kernel': ('linear', 'rbf', 'poly', 'rbf', 'sigmoid')})

print("tuned hpyerparameters :(best parameters) ",svm_cv.best_params_)
print("accuracy :",svm_cv.best_score_)

tuned hpyerparameters :(best parameters)  {'C': 1.0, 'gamma': 0.03162277660168379, 'kernel': 'sigmoid'}
accuracy : 0.8482142857142856
```

Predictive Analysis (Classification)

- Decision Tree

```
parameters = {'criterion': ['gini', 'entropy'],
              'splitter': ['best', 'random'],
              'max_depth': [2*n for n in range(1,10)],
              'max_features': ['auto', 'sqrt'],
              'min_samples_leaf': [1, 2, 4],
              'min_samples_split': [2, 5, 10]}

tree = DecisionTreeClassifier()
```

```
tree_cv = GridSearchCV(tree, parameters, cv=10)
tree_cv.fit(X_train, Y_train)
```

```
print("tuned hpyerparameters :(best parameters) ",tree_cv.best_params_)
print("accuracy :",tree_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters)  {'criterion': 'gini', 'max_depth': 4, 'max_features': 'auto', 'min_samples_leaf': 4, 'min_samples_split': 2, 'splitter': 'best'}
accuracy : 0.8892857142857142
```

Predictive Analysis (Classification)

- KNN

```
parameters = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],  
              'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],  
              'p': [1,2]}  
  
KNN = KNeighborsClassifier()
```

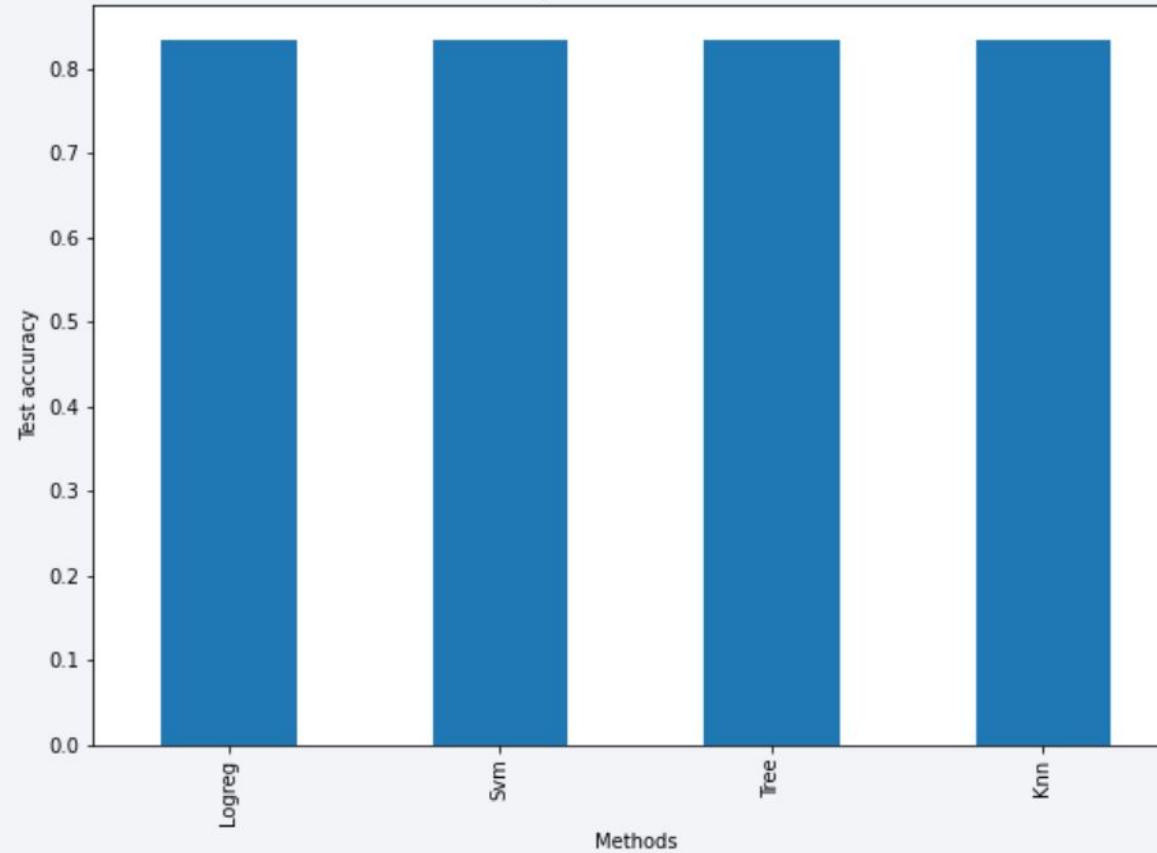
```
knn_cv = GridSearchCV(KNN, parameters, cv=10)  
knn_cv.fit(X_train, Y_train)  
  
GridSearchCV(cv=10, estimator=KNeighborsClassifier(),  
            param_grid={'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],  
                        'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],  
                        'p': [1, 2]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
print("tuned hpyerparameters :(best parameters) ",knn_cv.best_params_)  
print("accuracy :",knn_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters)  {'algorithm': 'auto', 'n_neighbors': 10, 'p': 1}  
accuracy : 0.8482142857142858
```

Classification Accuracy



Explanation:

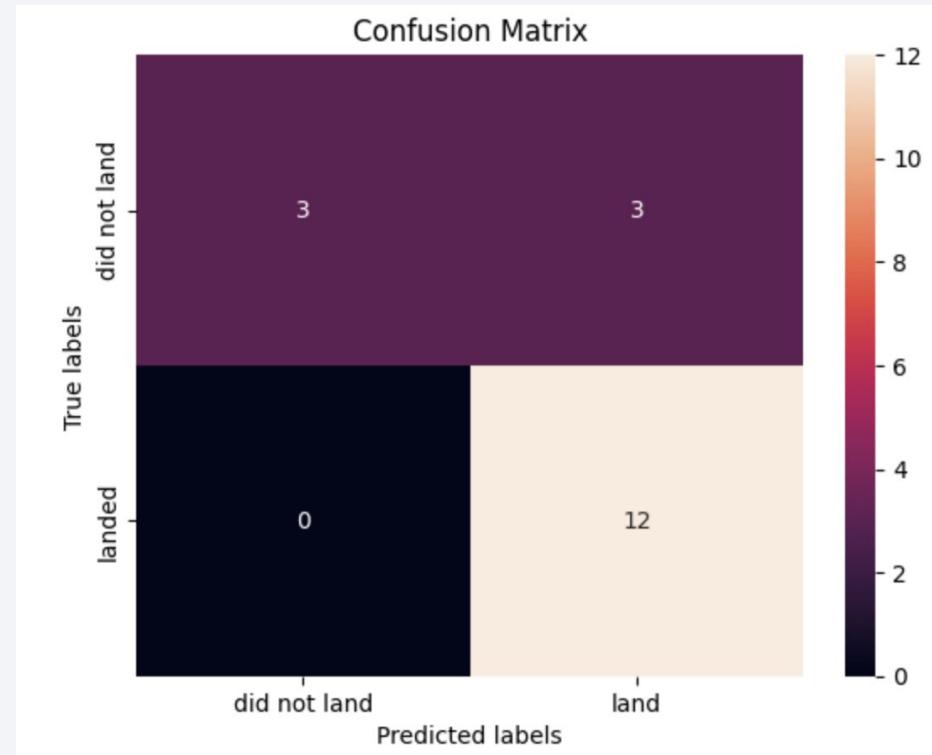
From the bar plot we can see that accuracy on test data for all methods are the same.

But on training data the best result has decision trees, its accuracy is 88.9%.

Decision Tree best params:

```
{'criterion': 'gini', 'max_depth': 4, 'max_features': 'auto', 'min_samples_leaf': 4, 'min_samples_split': 2, 'splitter': 'best'}  
accuracy : 0.8892857142857142
```

Confusion Matrix

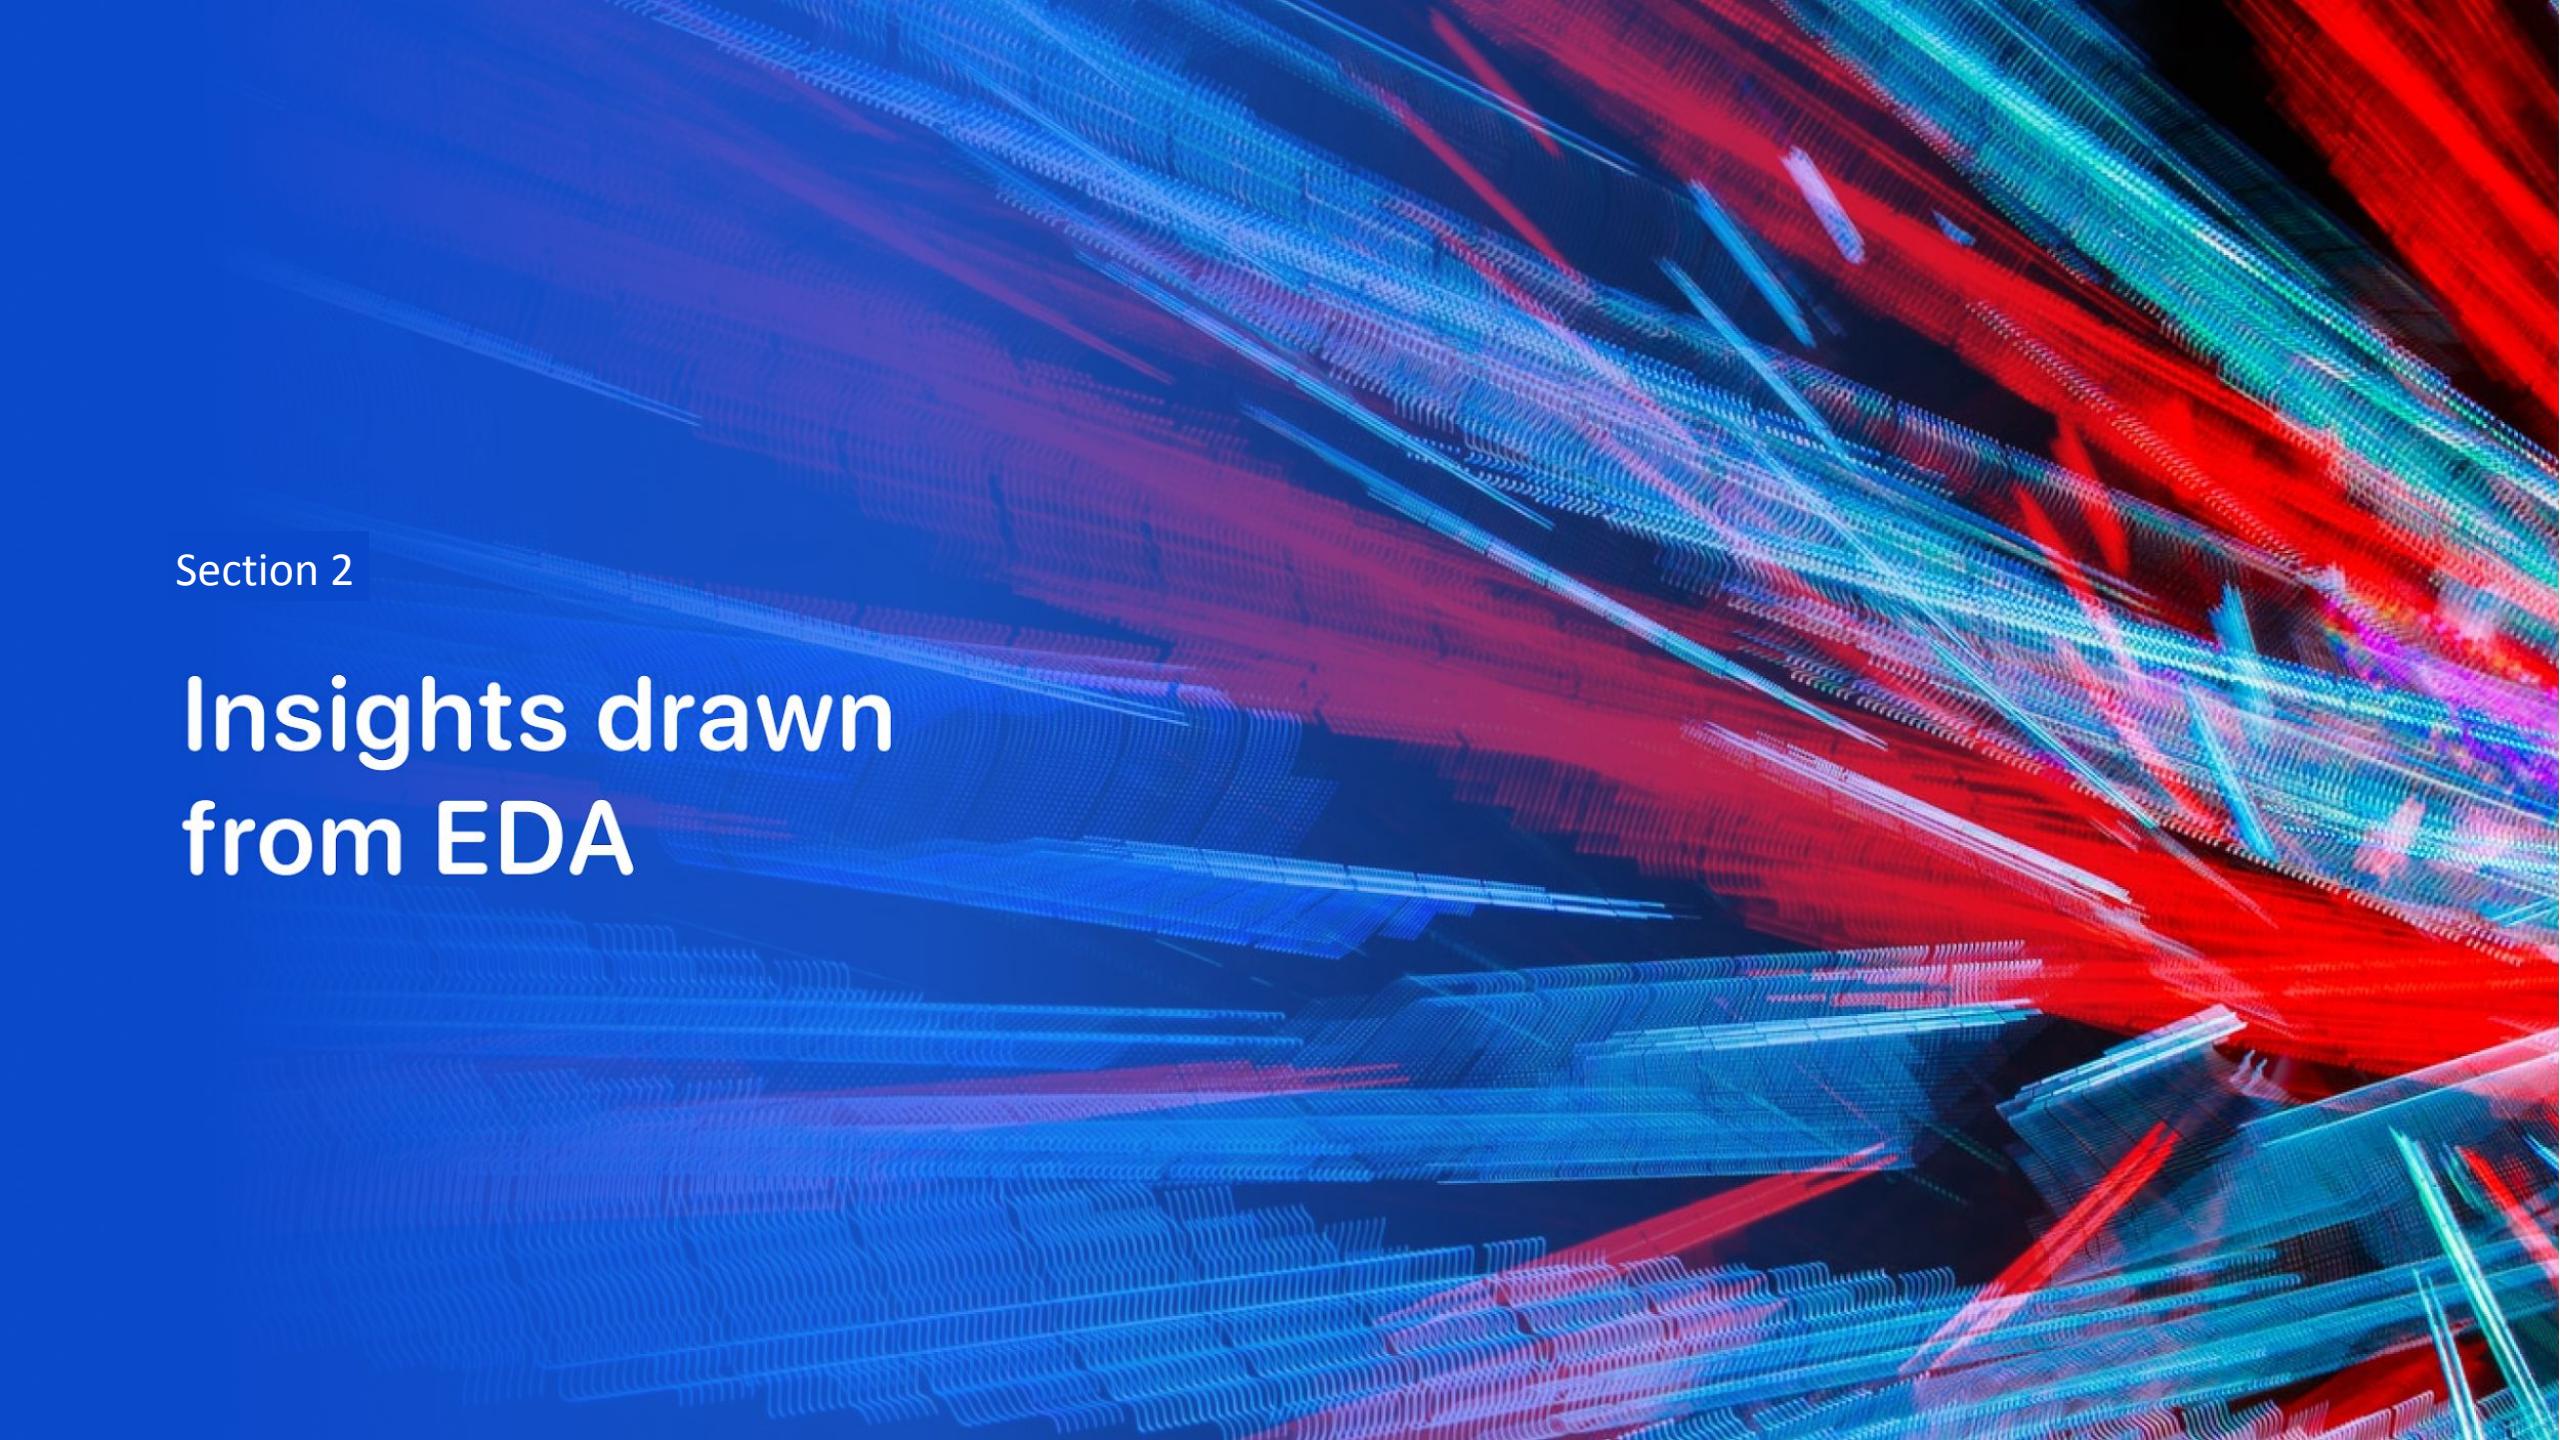


Explanation:

Confusion Matrix plots look the same for all methods. We can see from that graph that the main problem is with false positives.

Results

- Exploratory data analysis results
- Interactive analytics demo in screenshots
- Predictive analysis results

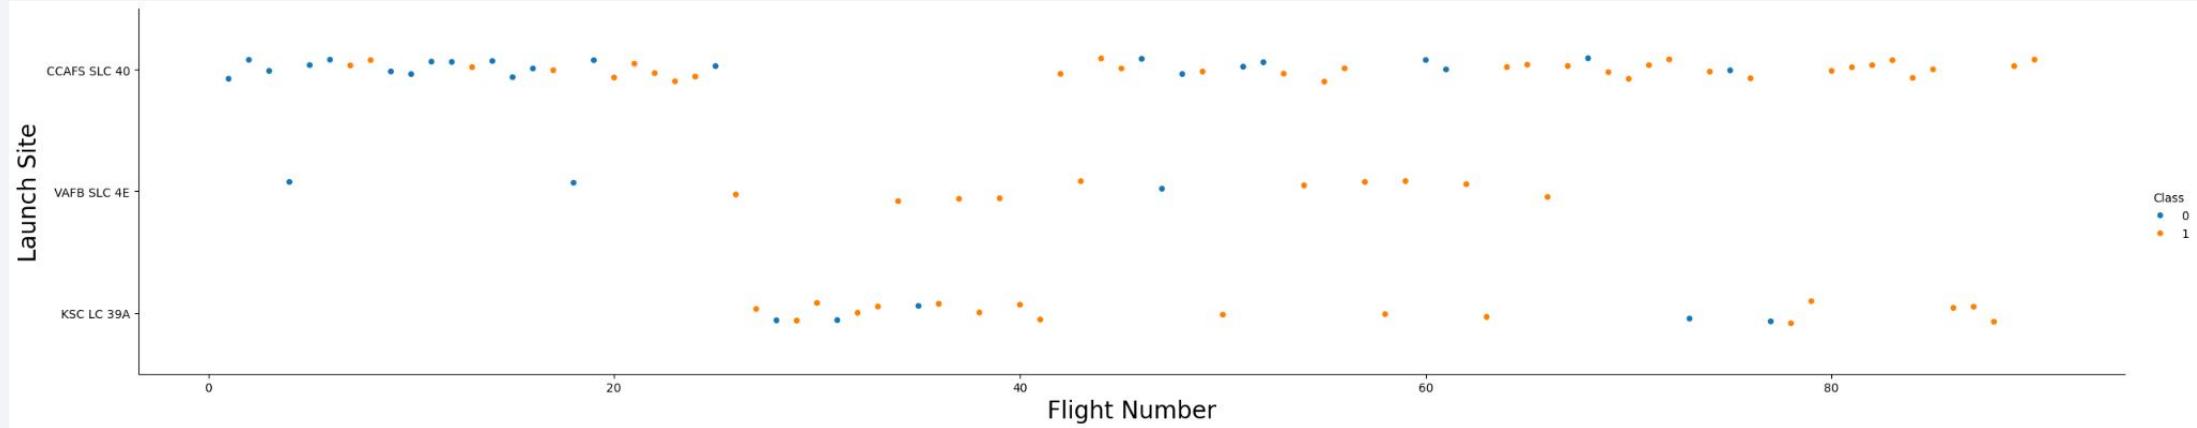
The background of the slide features a complex, abstract pattern of glowing lines. These lines are primarily blue and red, creating a sense of depth and motion. They appear to be composed of numerous small, glowing particles or dots, giving them a textured, almost liquid-like appearance. The lines converge and diverge, forming various shapes and directions across the dark, solid-colored background.

Section 2

Insights drawn from EDA

Flight Number vs. Launch Site

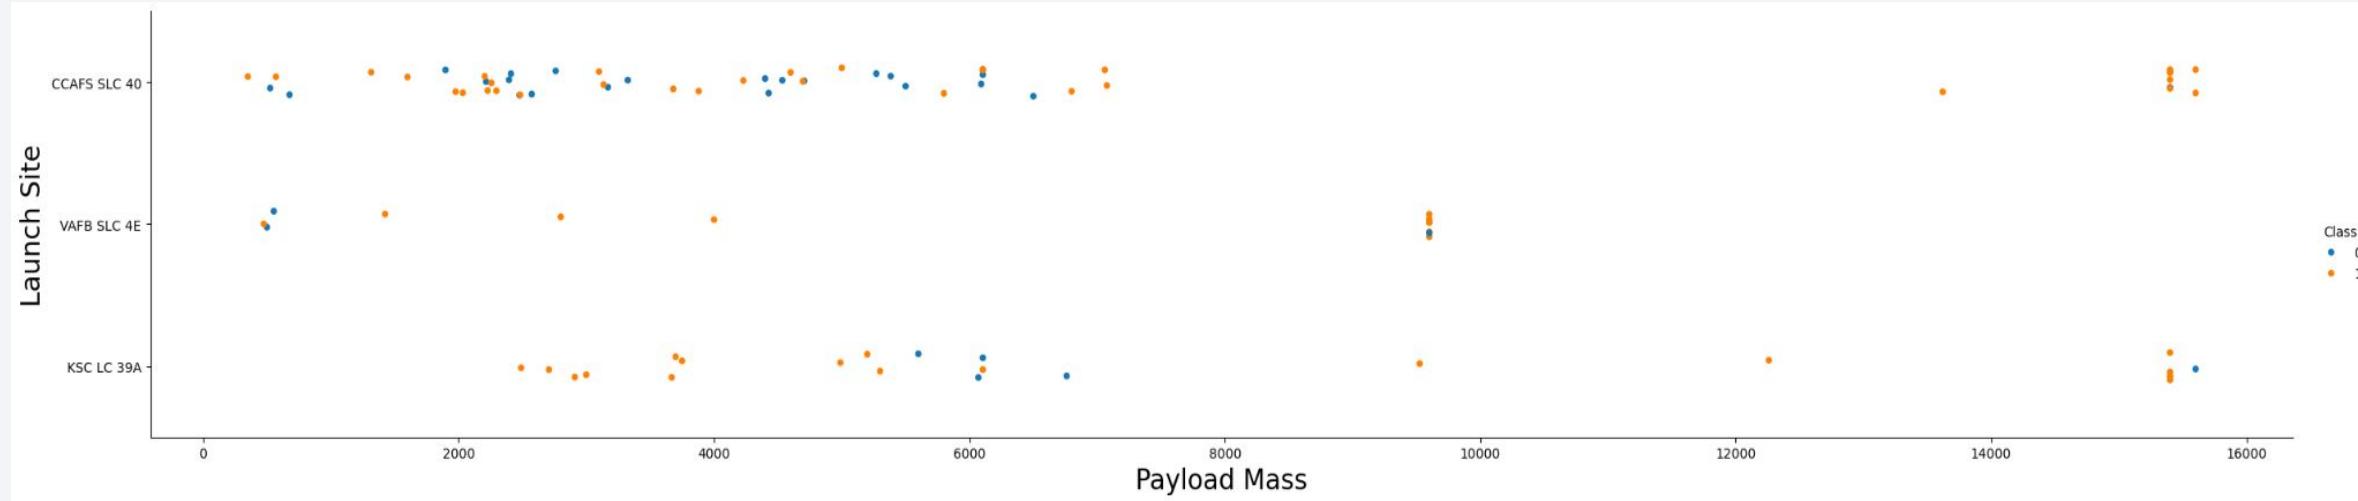
- Show a scatter plot of Flight Number vs. Launch Site



- We observe that with bigger FlightNumber success rate also increases.

Payload vs. Launch Site

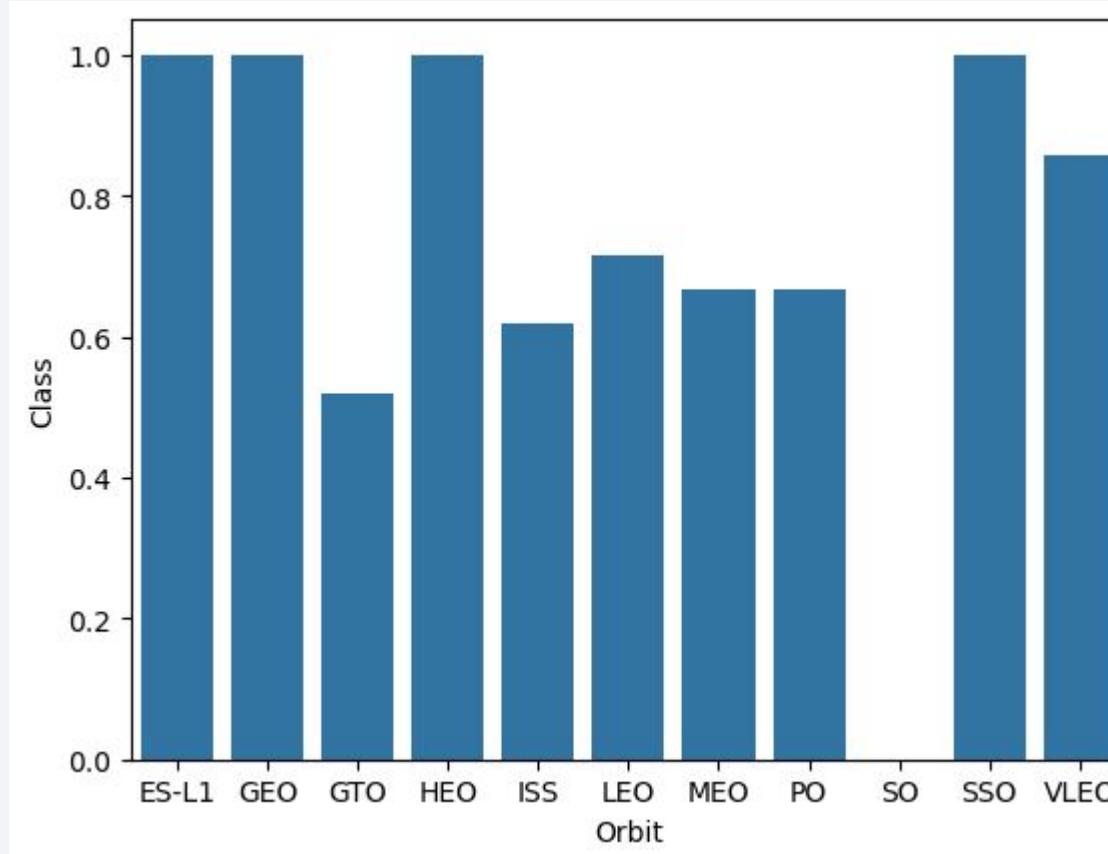
- Show a scatter plot of Payload vs. Launch Site



- Heavier payload may lead to successful landing but if it's too heavy it leads to failure.

Success Rate vs. Orbit Type

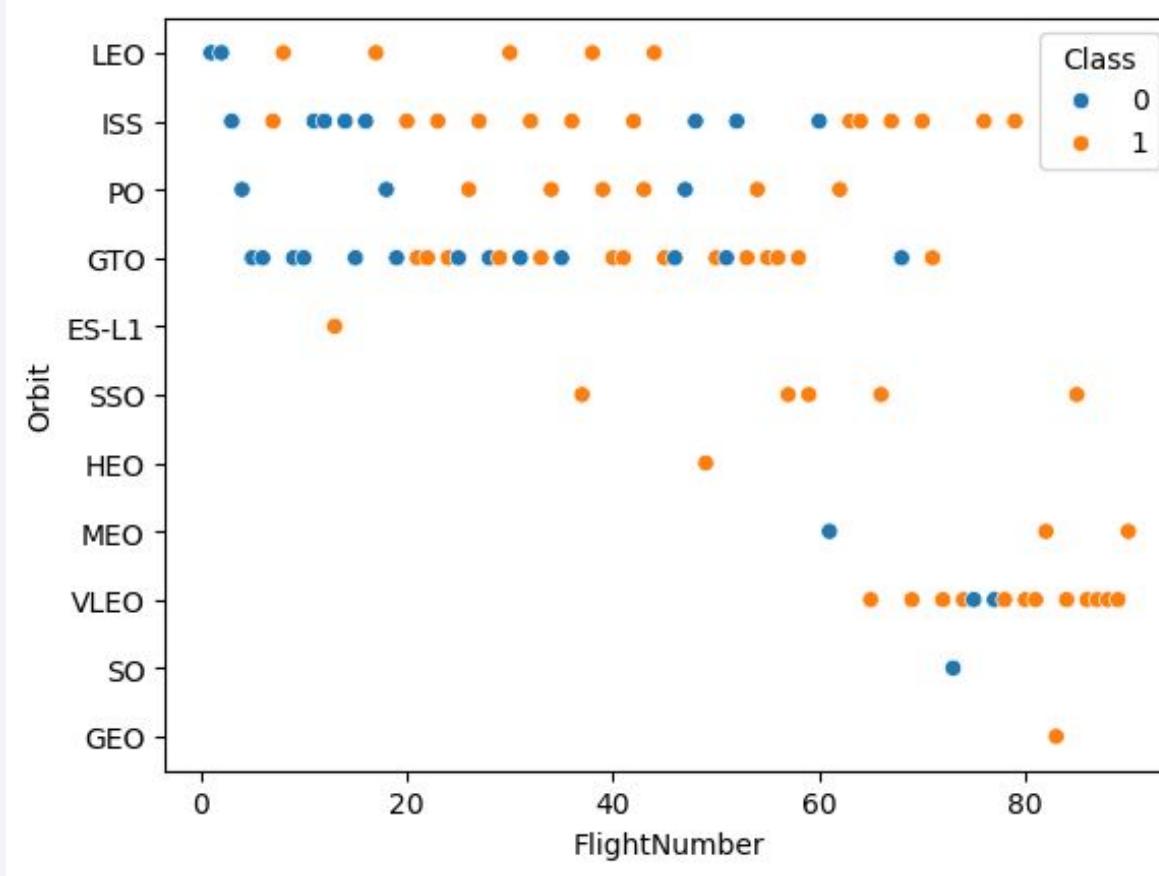
- Show a bar chart for the success rate of each orbit type



From the bar chart we can see that ES-L1, GEO, HEO, SSO have the best success rate.

Flight Number vs. Orbit Type

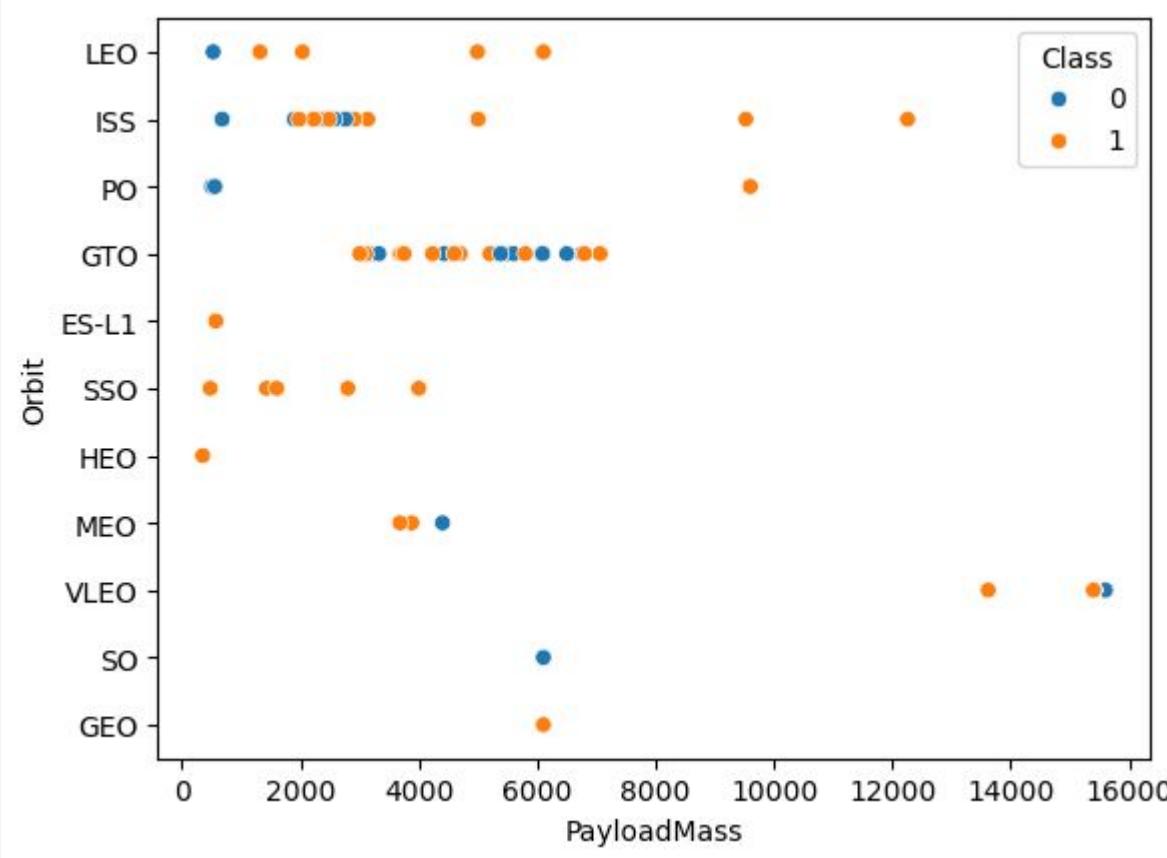
- Show a scatter point of Flight number vs. Orbit type



Success rate increases with the greater flight number but there are such orbits as GTO for example, where no relations with flight number found.

Payload vs. Orbit Type

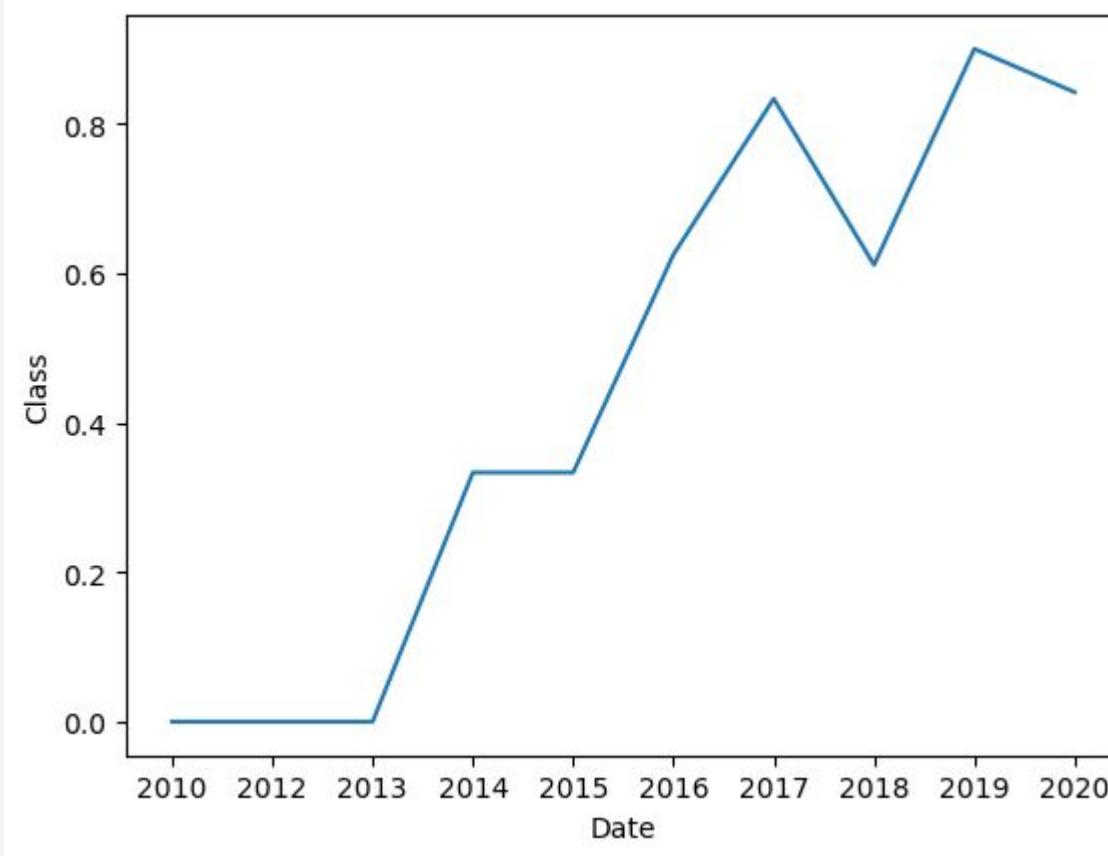
- Show a scatter point of payload vs. orbit type



Payload has great correlation with certain Orbit types. For example LEO, ISS have greater success rate with greater payload mass but for GTO, MEO it doesn't work.

Launch Success Yearly Trend

- Show a line chart of yearly average success rate



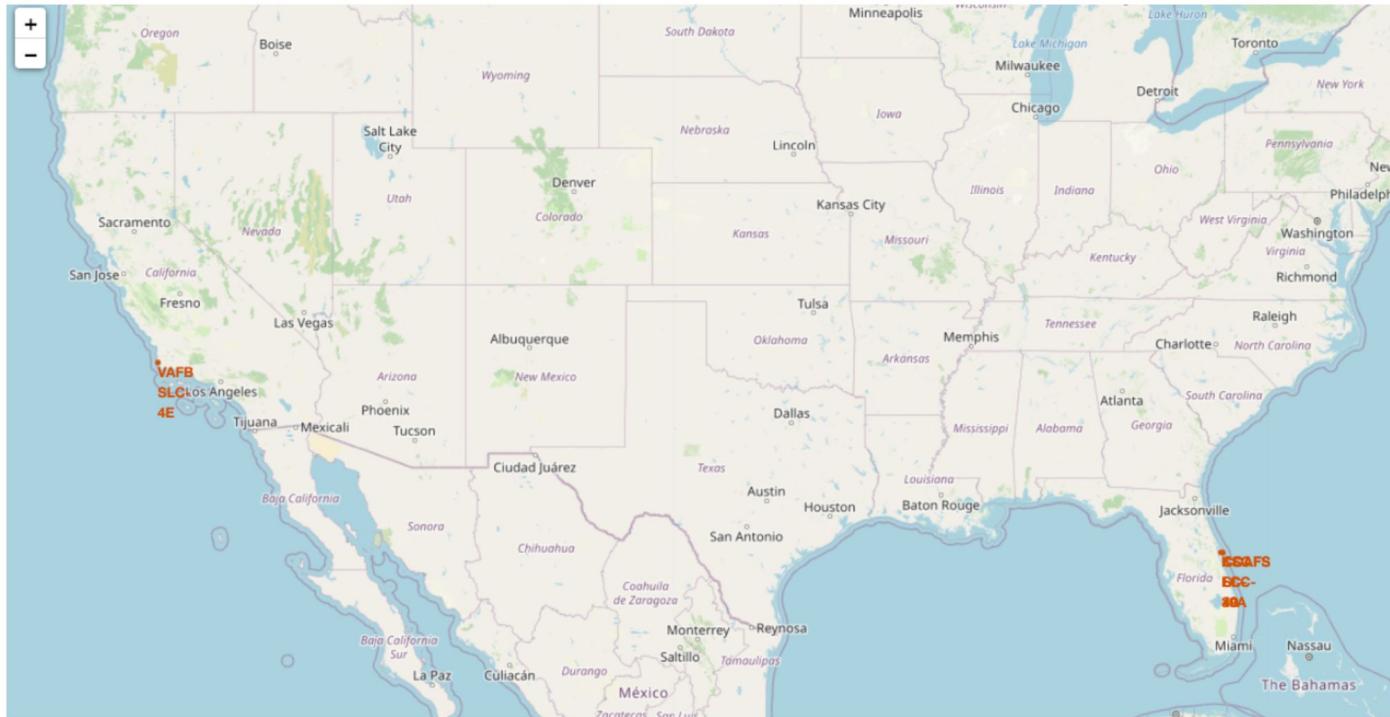
Since 2013 success rate keeps increasing with little drops.

The background of the slide is a photograph taken from space at night. It shows the curvature of the Earth's horizon against a dark blue sky. City lights are visible as small white dots, with larger clusters of lights indicating major urban areas. In the upper right corner, there is a faint, greenish glow of the aurora borealis or a similar atmospheric phenomenon.

Section 3

Launch Sites Proximities Analysis

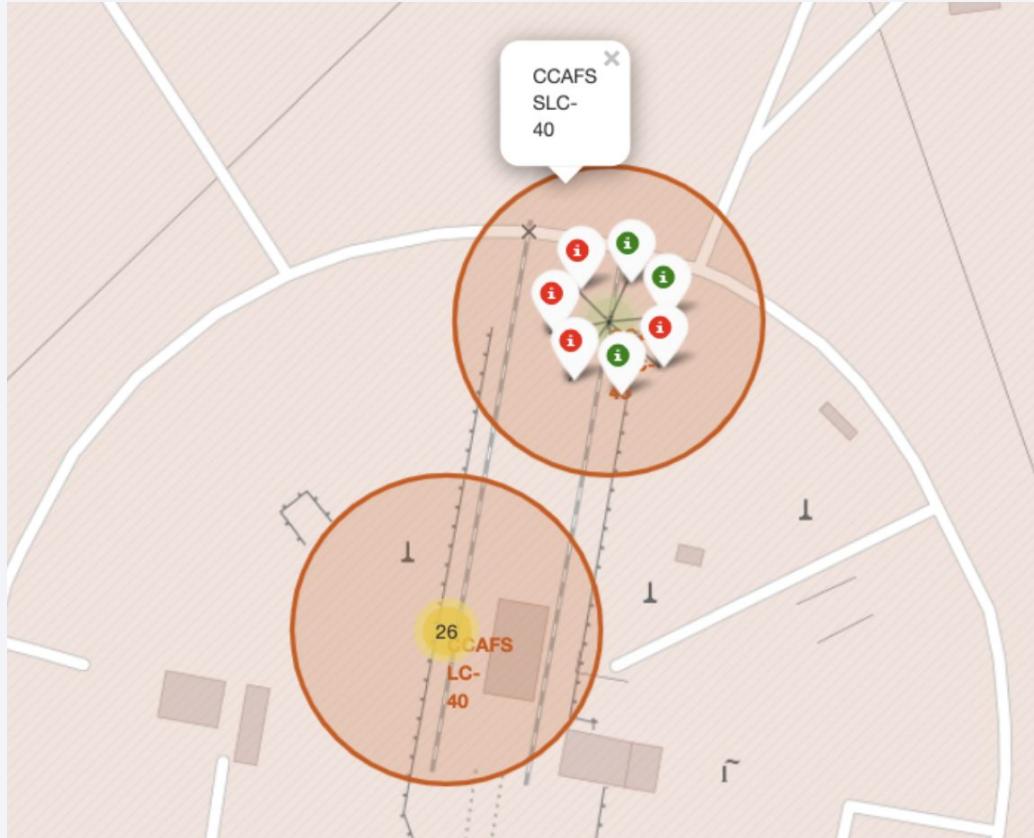
Folium Map - Ground stations



Explanation:

On that map we see that SpaceX launch sites located on the coast of US.

Folium Map - Marker clusters and labels

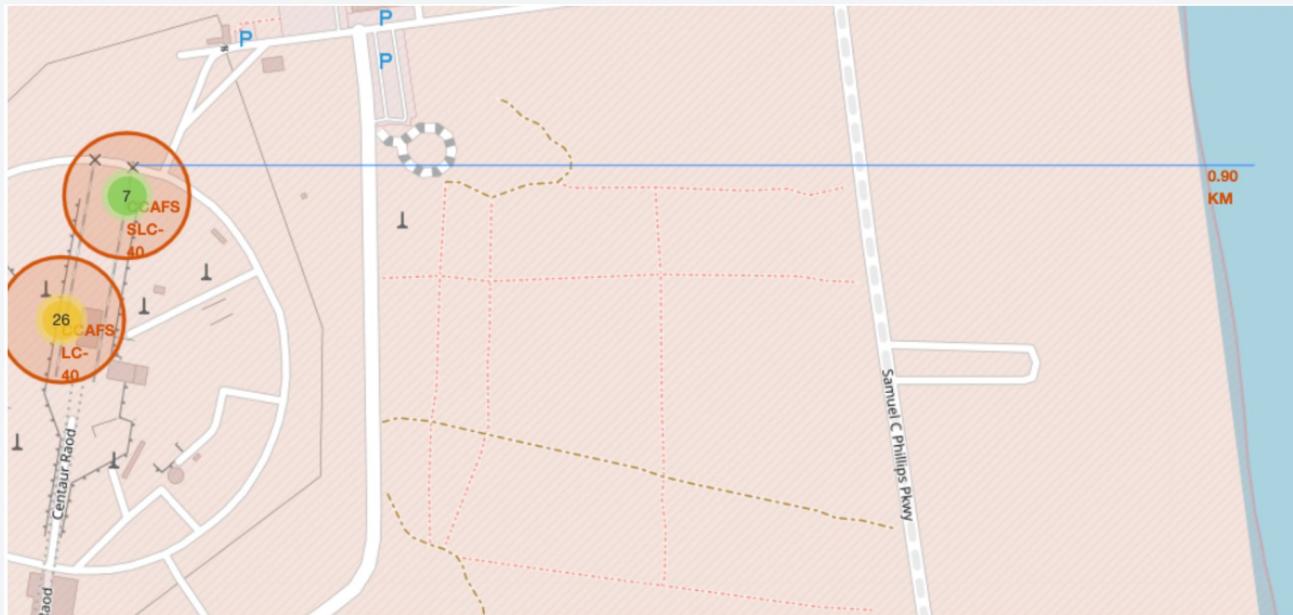


Explanation:

Red labels represent Failed launches and green labels represent Successful launches.

The most successful launch site is KSC LC-39A.

Folium Map - Distances between CCAFS SLC-40 and its proximities

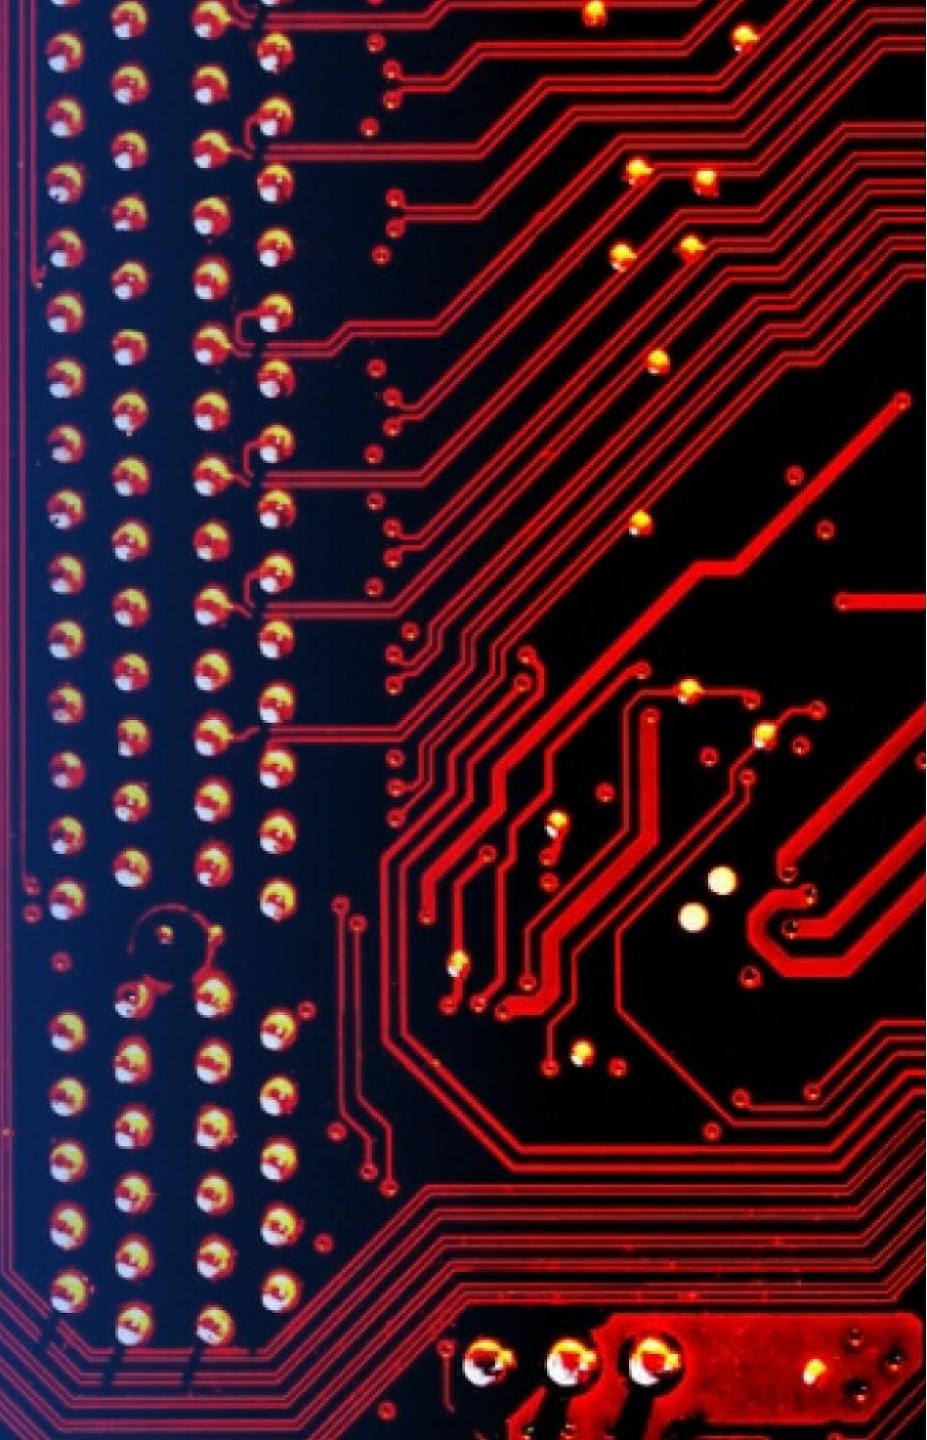


Explanation:

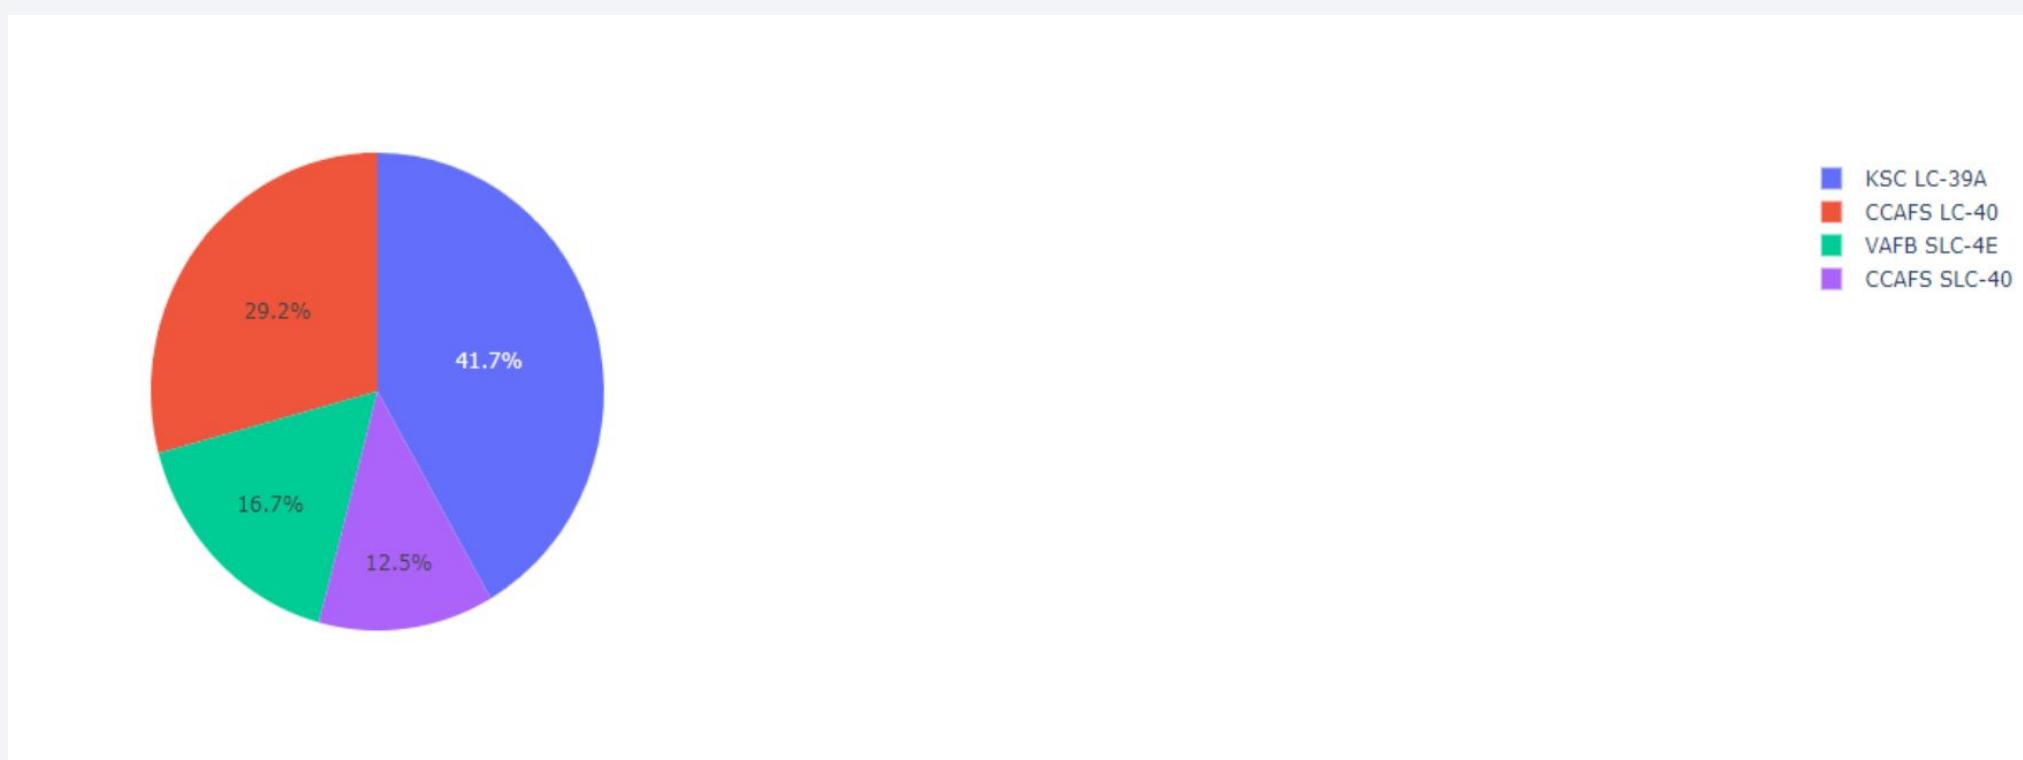
CCAFS SLC-40 is pretty close to all of its proximities such as railways, highways, coastline but it doesn't keep distance away from cities.

Section 4

Build a Dashboard with Plotly Dash



Plotly Dash - Total success rate



Explanation:

KSL LC-39A has the highest success rate.

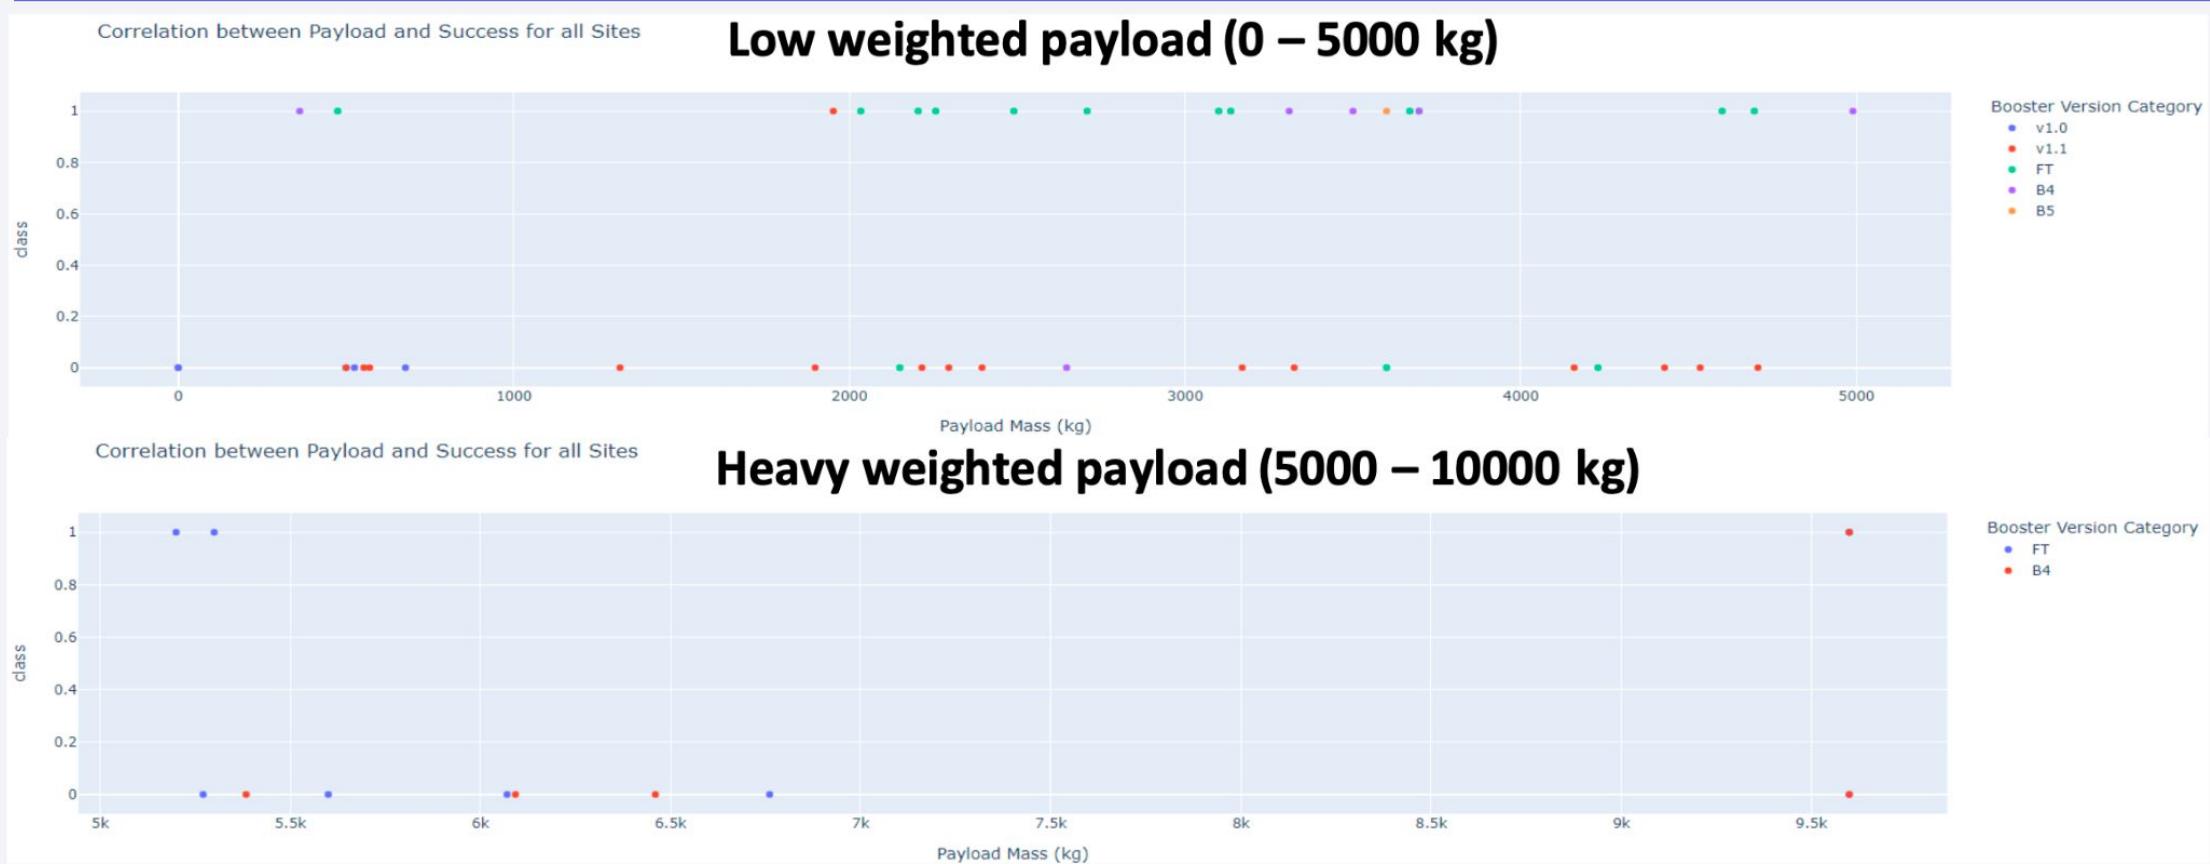
Plotly Dash - KSL LC-39A success rate



Explanation:

The launch site has 76.9% success launches and only 23.1% failed.

Plotly Dash - Correlation between Payload Mass vs Outcome for all launch sites



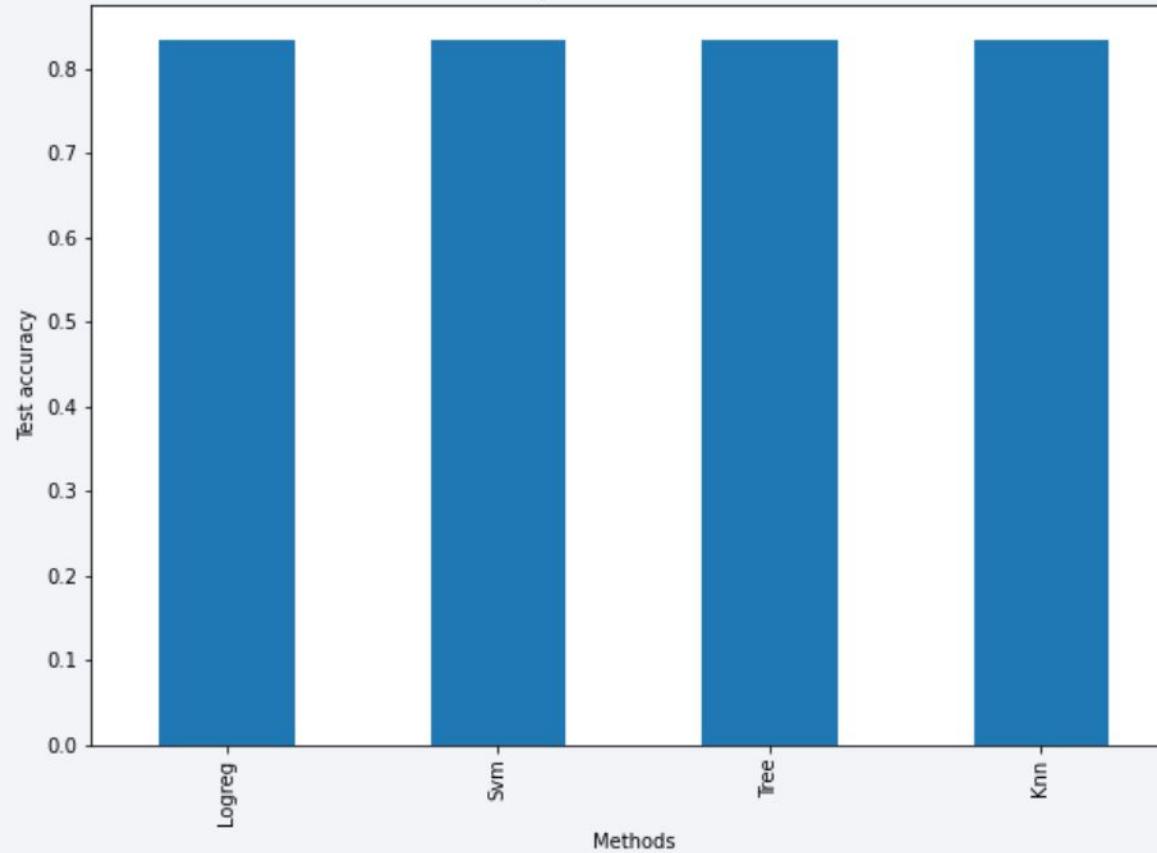
Explanation:

By comparing the plots we can see that success outcome for specific Booster versions is highest payload in [2000, 5000] kg.

Section 5

Predictive Analysis (Classification)

Classification Accuracy



Explanation:

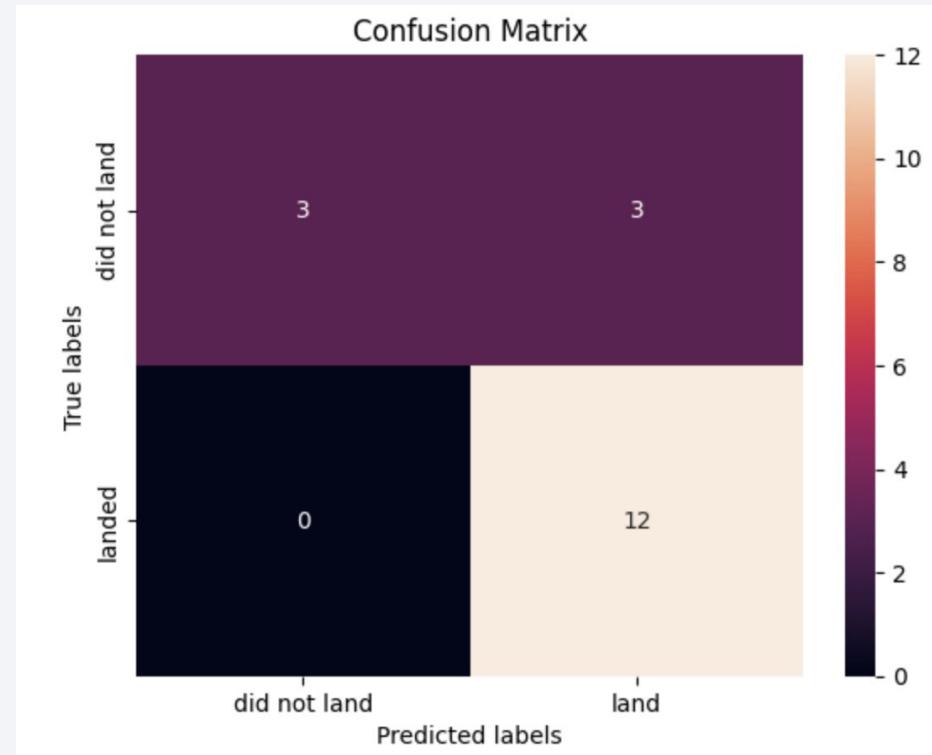
From the bar plot we can see that accuracy on test data for all methods are the same.

But on training data the best result has decision trees, its accuracy is 88.9%.

Decision Tree best params:

```
{'criterion': 'gini', 'max_depth': 4, 'max_features': 'auto', 'min_samples_leaf': 4, 'min_samples_split': 2, 'splitter': 'best'}  
accuracy : 0.8892857142857142
```

Confusion Matrix



Explanation:

Confusion Matrix plots look the same for all methods. We can see from that graph that the main problem is with false positives.

Conclusions

- We can see that success outcome correlates with such variables as flight number, launch site, the orbit.
- The best success rates we can see on the following orbits: GEO, HEO, SSO, ES-L1.
- The payload mass highly correlated with successful outcome, payloads mass in such range as [2000, 5000] kg leads to successful landing but it's true only for specific orbits.
- We cannot explain correlation between launch site and successful landing with the provided data.
- The best result of prediction corresponds to Decision Tree, even when on test dataset we see the same results among methods on training data Decision Tree shows the best accuracy 88.9%.

Appendix

- https://github.com/alexkarpovich/datascience-coursera/blob/main/IBM-DS0321EN-SkillsNetwork_labs_module_2_jupyter-labs-eda-dataviz.ipynb.jupyterlite.ipynb
- https://github.com/alexkarpovich/datascience-coursera/blob/main/IBM-DS0321EN-SkillsNetwork_labs_module_3_lab_jupyter_launch_site_location.jupyterlite.ipynb
- https://github.com/alexkarpovich/datascience-coursera/blob/main/IBM-DS0321EN-SkillsNetwork_labs_module_4_SpaceX_Machine_Learning_Prediction_Part_5.jupyterlite.ipynb
- [https://github.com/alexkarpovich/datascience-coursera/blob/main/jupyter-labs-eda-sql-coursera_sqlite%20\(1\).ipynb](https://github.com/alexkarpovich/datascience-coursera/blob/main/jupyter-labs-eda-sql-coursera_sqlite%20(1).ipynb)
- <https://github.com/alexkarpovich/datascience-coursera/blob/main/jupyter-labs-spacex-data-collection-api.ipynb>
- <https://github.com/alexkarpovich/datascience-coursera/blob/main/jupyter-labs-webscraping.ipynb>
- [https://github.com/alexkarpovich/datascience-coursera/blob/main/labs-jupyter-spacex-Data%20wrangling%20\(1\).ipynb](https://github.com/alexkarpovich/datascience-coursera/blob/main/labs-jupyter-spacex-Data%20wrangling%20(1).ipynb)
- https://github.com/alexkarpovich/datascience-coursera/blob/main/spacex_dash_app.py

Thank you!

